

# Drop-in text for Sec. VI-E (deviation bound) and Sec. VI-B (ramp tracking)

numbering continues from Eq. (86); symbols follow the manuscript

## VI-E continuation: impulse response and the bound on $e_\varphi$

After Proposition 1 / Eq. (86).

The homogeneous part of (86) is spanned by  $\cosh C_2\tau$  and  $\sinh C_2\tau$  with  $C_2 = \sqrt{Wz_{CoM}/J_P}$  from (20). Its impulse response — the solution of  $J_P\ddot{h} - Wz_{CoM}h = \delta(\tau)$  from rest — follows from the momentum jump  $J_P[\dot{h}(0^+) - \dot{h}(0^-)] = 1$ , the position being unable to jump:

$$h_\varphi(\tau) = \frac{\sinh(C_2\tau)}{J_PC_2}, \quad h_\varphi(0) = 0, \quad \dot{h}_\varphi(0) = \frac{1}{J_P}. \quad (87)$$

Because  $h_\varphi(0) = 0$ , the boundary term generated by differentiating a convolution with a variable upper limit,

$$\frac{d}{d\tau} \int_0^\tau h_\varphi(\tau-s)\rho(s)ds = \underbrace{h_\varphi(0)}_{=0} \rho(\tau) + \int_0^\tau \dot{h}_\varphi(\tau-s)\rho(s)ds,$$

vanishes for *any* forcing. The zero initial conditions in (86) therefore set both homogeneous coefficients to zero — no assumption on  $\rho(0)$  is required — and the deviation is the Duhamel integral alone,

$$e_\varphi(\tau) = \int_0^\tau h_\varphi(\tau-s)\rho(s)ds, \quad e_\omega(\tau) = \int_0^\tau \dot{h}_\varphi(\tau-s)\rho(s)ds. \quad (88)$$

**The two channels.** Two terms of (79) leave the estimator span: the gravity remainder (83) and the bilinear ground-effect coupling (79)(iv),

$$\rho(\tau) = \rho_\varphi(\tau) + \rho_{GE}(\tau), \quad \rho_{GE}(\tau) = \beta_M \dot{M}\tau\varphi. \quad (89)$$

Both vanish at the onset — with  $\varphi_{\text{nom}} \sim \tau^3$  the gravity remainder is  $\mathcal{O}(\varphi^2) \sim \tau^6$  and the bilinear term is  $\tau \cdot \tau^3 = \tau^4$  — so each is largest at the window end.

**An envelope bound that can be verified by hand.** The estimator consumes the *rate*: the two-segment fit is applied to  $\omega = C_1(\cosh C_2\tau - 1)$  reconstructed from the gyro, and the onset is the argument that minimises its residual. The quantity that must be bounded is therefore  $e_\omega$ , and  $e_\varphi$  enters only through the gravity remainder that generates  $\rho$ . Replacing  $\rho$  by a constant  $\bar{\rho}$  and integrating the two kernels,  $\int_0^\tau \dot{h}_\varphi = \sinh(C_2\tau)/(J_PC_2)$  and  $\int_0^\tau h_\varphi = (\cosh C_2\tau - 1)/(J_PC_2^2)$ , with  $J_PC_2 = Wz_{CoM}/C_2$  and  $J_PC_2^2 = Wz_{CoM}$ ,

$$|e_\omega(\tau_{\text{end}})| \leq \frac{\bar{\rho}C_2}{Wz_{CoM}} \sinh x, \quad |e_\varphi(\tau_{\text{end}})| \leq \frac{\bar{\rho}}{Wz_{CoM}} (\cosh x - 1), \quad x \triangleq C_2\tau_{\text{end}}. \quad (90)$$

The inertia cancels in both. Only  $Wz_{CoM}$ ,  $C_2$ , the dimensionless group  $x$  and the forcing level enter.

**The window length is bounded a priori by the tilt cap.** Neither channel maximum nor the kernel factor requires  $x$  to be solved for. With the nominal excursion of (20),

$$\varphi_{\text{end}} = \frac{\dot{M}}{Wz_{CoM}C_2} \Lambda(x), \quad \Lambda(x) \triangleq \sinh x - x = \sum_{k \geq 1} \frac{x^{2k+1}}{(2k+1)!}. \quad (91)$$

Every term of  $\Lambda$  is positive, so  $\Lambda(x) \geq x^3/6$ , and since  $\Lambda$  increases, imposing the tilt cap  $\varphi_{\text{end}} \leq \varphi_{\text{max}}$  inverts in closed form:

$$x \leq \bar{x} \triangleq \left( \frac{6\varphi_{\text{max}}Wz_{CoM}C_2}{\dot{M}} \right)^{1/3}, \quad \tau_{\text{end}} \leq \frac{\bar{x}}{C_2} = \left( \frac{6\varphi_{\text{max}}Wz_{CoM}}{C_2^2\dot{M}} \right)^{1/3}, \quad (92)$$

the second form using  $J_P = Wz_{CoM}/C_2^2$ . The retained term is the inertia-only rise  $\varphi = \dot{M}\tau^3/(6J_P)$  obtained by dropping the destabilising  $Wz_{CoM}\varphi$ : it is the slowest admissible ascent and therefore reaches the cap latest, which is why (92) bounds  $x$  from above and why the worst case is the *slowest* ramp, not the fastest. The moment increment over the window and the coupling channel follow with no fitted quantity at all,

$$\begin{aligned}\Delta M_{\text{win}} &= \dot{M}\tau_{\text{end}} \leq (6\varphi_{\text{max}}J_P\dot{M}^2)^{1/3}, \\ \rho_{GE,\text{max}} &= |\beta_M|\Delta M_{\text{win}}\varphi_{\text{end}} \leq |\beta_M|\varphi_{\text{max}}(6\varphi_{\text{max}}J_P\dot{M}^2)^{1/3}.\end{aligned}\tag{93}$$

and the kernel factors of (90) close on the same constant through the identity  $\cosh x - 1 = \Lambda(x) + (x - 1 + e^{-x})$  with  $0 \leq x - 1 + e^{-x} < x$ , together with  $\sinh x = \Lambda(x) + x$ :

$$\cosh x - 1 \leq \sinh x \leq \Lambda(x) + \bar{x} \leq \frac{\bar{x}^3}{6} + \bar{x}.\tag{94}$$

Both kernels therefore share one bracket.<sup>1</sup>

**Why the envelope alone does not settle the question.** Nothing beyond  $|\rho| \leq \rho_{\text{max}}$  enters (90): no monotonicity, no shape, no property of the trajectory but the window length. That generality is also its cost. A constant envelope spreads the forcing over the whole window, whereas both channels are concentrated at its end ( $\tau^6$  and  $\tau^4$ ), and the mismatch grows with  $x$ , i.e. with the *slowest* ramps. Evaluated with  $\bar{\rho} = \rho_{\text{max}}$  over the box, (90) returns a relative deviation of 446% at  $\dot{M} = 0.10 \text{ N m/s}$  — larger than the signal it bounds. The envelope form is a sanity check on the algebra, not a result.

What restores it is the one structural fact available for free. Both channels rise monotonically from zero at the onset, while the kernel seen as a function of the integration variable,  $s \mapsto \dot{h}_\varphi(\tau_{\text{end}} - s)$ , *decreases* — the reflection  $s \mapsto \tau_{\text{end}} - s$  reversing a cosh that increases in its own argument. Chebyshev’s integral inequality for factors of *opposite* monotonicity therefore applies, in the direction needed:  $\int_a^b fg \leq (b - a)^{-1} \int_a^b f \int_a^b g$ , which follows by integrating  $(f(u) - f(v))(g(u) - g(v)) \leq 0$  over  $[a, b]^2$ . Applied to (88) it replaces each maximum in (90) by the channel’s time average over the window. With the nominal shapes  $\rho_\varphi \propto \Lambda(u)^2$  and  $\rho_{GE} \propto u \Lambda(u)$ ,  $u = C_2\tau$ , the amplitudes cancel in the ratio and

$$\begin{aligned}\bar{\rho} &= R_\varphi(x)\rho_{\varphi,\text{max}} + R_{GE}(x)\rho_{GE,\text{max}}, \\ R_\varphi(x) &= \frac{1}{x} \frac{\int_0^x \Lambda(u)^2 du}{\Lambda(x)^2} \leq \frac{1}{7}, \quad R_{GE}(x) = \frac{1}{x} \frac{\int_0^x u \Lambda(u) du}{x \Lambda(x)} \leq \frac{1}{5}.\end{aligned}\tag{95}$$

The two constants are simply the orders of the two channels at the onset: a signal going as  $\tau^n$  averages to  $1/(n+1)$  of its endpoint value, and the channels are sixth and fourth order. That these are the suprema follows from rescaling the window,  $u = xt$ , which turns either ratio into  $R(x) = \int_0^1 f(xt)/f(x) dt$ ; the logarithmic derivative of that integrand is  $[n_{\text{eff}}(xt) - n_{\text{eff}}(x)]/x$  with  $n_{\text{eff}} = uf'/f$  the local order, so  $R$  decreases wherever  $n_{\text{eff}}$  rises. It does in both cases —  $n_{\text{eff}}$  equals  $2\Psi_\varphi$  and  $1 + \Psi_\varphi$  respectively, with  $\Psi_\varphi(u) = u\Lambda'/\Lambda$  increasing from 3 — because  $\Lambda$  stiffens from cubic to exponential and concentrates each channel ever more sharply at the end of the window. Over the flown  $x \in [1.79, 5.20]$  the factors are 0.089–0.133 and 0.145–0.191, so (95) gives away 10–60%; the same evaluation that returned 446% returns 41.8% with it in place.

**Relative form: the result, free of  $x$ .** Dividing the rate bound of (90) by the nominal rate the estimator fits,  $\omega_{\text{nom, end}} = C_1(\cosh x - 1)$  with  $C_1 = \dot{M}/Wz_{CoM}$ , and writing  $\Delta M_{\text{win}} = \dot{M}\tau_{\text{end}} = \dot{M}x/C_2$  for the moment increment over the window,

$$\frac{|e_\omega|_{\text{end}}}{\omega_{\text{nom, end}}} \leq \frac{\Psi(x)\bar{\rho}}{\Delta M_{\text{win}}}, \quad \Psi(x) = \frac{x \sinh x}{\cosh x - 1} = x \coth \frac{x}{2},\tag{96}$$

the closed form following from  $\cosh x - 1 = 2 \sinh^2(x/2)$  and  $\sinh x = 2 \sinh(x/2) \cosh(x/2)$ ;  $\Psi \geq 2$  everywhere, which is the inequality the exponent-perturbation argument of Sec. VI-D already uses. Here the two approximations work against each other:  $\Psi$  grows with  $x$  while the averages of (95) shrink with it, and each product is bounded by its  $x \rightarrow \infty$  limit,  $\frac{1}{2}$  and 1. Both bounds are exact rather than sampled. Substituting (95) and  $\Psi = x \coth(x/2)$ , the two claims read

$$\int_0^x \Lambda(u)^2 du \leq \frac{1}{2} \tanh \frac{x}{2} \Lambda(x)^2, \quad \int_0^x u \Lambda(u) du \leq x \tanh \frac{x}{2} \Lambda(x),\tag{99}$$

<sup>1</sup>The step matters.  $\Lambda(x)$  is known as an *equality* — the tilt cap supplies it — so only the residual linear term needs  $\bar{x}$ . Substituting  $\bar{x}$  into  $\sinh$  instead is exponentially lossy: at the slowest ramp  $\bar{x} = 7.94$  gives  $\sinh \bar{x} = 1404$  against  $\Lambda + \bar{x} = 91.4$ , a factor of fifteen thrown away for nothing.

with  $\Lambda(u) = \sinh u - u$ ; the half-angle tangent enters because  $\tanh(x/2) = \Lambda'(x)/\sinh x$ . Each step from (95)–(96) to (99) divides by a quantity that is positive for  $x > 0$  —  $\coth(x/2)$ ,  $\Lambda$ ,  $\Lambda^2$  and  $\cosh^2(x/2)$  — so (99) is equivalent to the two claims and not merely sufficient for them.

Both are proved from one corollary of the mean value theorem, used three times and nowhere supplemented. *If  $f$  is continuous on  $[0, b]$ , differentiable on  $(0, b)$  with  $f' \geq 0$ , and  $f(0) = 0$ , then  $f \geq 0$  on  $[0, b]$ :* for  $t \in (0, b]$  the theorem gives  $f(t) = f(t) - f(0) = f'(\xi)t \geq 0$  for some  $\xi \in (0, t)$ . Write the first claim of (99) as  $D(x) \triangleq \frac{1}{2} \tanh(x/2) \Lambda(x)^2 - \int_0^x \Lambda^2$ , so  $D(0) = 0$ ; differentiating and factoring out  $\Lambda > 0$  leaves  $D' = \Lambda B$  with  $B = \phi/(4 \cosh^2(x/2))$  and

$$\phi(x) \triangleq x + 2x \cosh x - 3 \sinh x, \quad \phi''(x) = \sinh x + 2x \cosh x \geq 0. \quad (100)$$

Applying the corollary to  $\phi'$  (whose derivative is  $\phi'' \geq 0$  and which vanishes at 0) gives  $\phi' \geq 0$ ; applying it to  $\phi$  gives  $\phi \geq 0$ , hence  $B \geq 0$  and  $D' = \Lambda B \geq 0$ ; applying it to  $D$  gives  $D \geq 0$ , which is the claim. The second inequality of (99) follows the same way; equivalently its power series has the general coefficient  $[2^{m-1}/m + (m^2 - 6m + 2)/3]/(m-1)!$  with  $m = 2n$ , zero at  $m = 2, 4$  and positive for  $m \geq 6$ . No monotonicity of the products is used, and the limits  $\frac{1}{2}$  and 1 enter the proof nowhere: they are values the constants are held to, and the  $x \rightarrow \infty$  limit shows only that neither can be lowered. Hence

$$\boxed{\frac{|e_\omega|_{\text{end}}}{\omega_{\text{nom, end}}} \leq \frac{\frac{1}{2}\rho_{\varphi, \text{max}} + \rho_{GE, \text{max}}}{\Delta M_{\text{win}}}, \quad |\Delta M_{\text{crit}}| \leq \frac{\rho_{\varphi, \text{max}}}{7} + \frac{\rho_{GE, \text{max}}}{5}} \quad (97)$$

independently of  $x$ ,  $C_2$ ,  $J_P$  and  $Wz_{CoM}$ : the window length that (92) went to some trouble to bound does not survive into the conclusion, and is needed only to bound  $\rho_{GE, \text{max}}$  through (93).

The second inequality of (97) is the same argument carried one step further, to the quantity actually reported. The onset minimises the two-segment residual, so perturbing the data by  $e_\omega$  displaces the minimiser by  $\delta = \langle e_\omega, \partial_{t_0} \hat{\omega} \rangle / \|\partial_{t_0} \hat{\omega}\|^2$  with  $\partial_{t_0} \hat{\omega} = -C_1 C_2 \sinh(C_2 \tau)$ ; multiplying by  $\dot{M}$  and substituting (88), then exchanging the order of integration, collapses the double integral to a weighted average of the forcing itself,

$$\Delta M_{\text{crit}} = - \int_0^{\tau_{\text{end}}} \rho(s) w(s) ds, \quad w(s) = \frac{Wz_{CoM} \int_s^{\tau_{\text{end}}} \sinh(C_2 \tau) \cosh(C_2(\tau - s)) d\tau}{J_P C_2 \|\sinh(C_2 \tau)\|^2}. \quad (98)$$

The weight is non-negative, being an integral of a product of hyperbolic functions of non-negative arguments; it is non-increasing in  $s$ , both the integration range and the cosh factor shrinking as  $s$  grows; and it integrates to unity exactly. The last is worth a line, since it is what lets  $\bar{\rho}$  enter (97) bare. Integrating (98) over  $s$  and exchanging the order — the region  $0 \leq s \leq \tau \leq \tau_{\text{end}}$  being the same triangle read either way — the inner integral is  $\int_0^\tau \cosh(C_2(\tau - s)) ds = \sinh(C_2 \tau)/C_2$ , which reproduces the norm in the denominator, so  $\int w = Wz_{CoM}/(J_P C_2^2) = 1$  by the same identity that removes the inertia elsewhere. Neither the window length nor any constant survives, and a uniform unmodelled moment therefore displaces the identified threshold one-for-one as a consequence rather than as an assumption. The first two properties alone would give  $|\Delta M_{\text{crit}}| \leq \rho_{\text{max}}$ . The third makes  $\rho$  and  $w$  oppositely monotone, so Chebyshev applies a second time and returns  $|\Delta M_{\text{crit}}| \leq \bar{\rho}$ , which with (95) is the stated bound.

The constants differ between the two halves of (97) because the two questions do. The first asks how far the fitted rate curve is displaced, so the kernel integral  $\Psi$  multiplies  $\bar{\rho}$  and the products climb to  $\frac{1}{2}$  and 1; the second asks how far the *minimiser* moves, and by (98) that is a normalised correlation with no kernel gain left over, so  $\bar{\rho}$  enters bare and the constants are the reduction factors themselves. The threshold bound is the stronger of the two by about a factor of four, which is why it, and not the relative rate bound, is what the offset budget inherits. A reviewer can evaluate either from numbers already in the paper.

**How the pivot inertia enters, and how little of it is needed.** The two bounds want  $J_P$  from opposite sides, and the parallel-axis identity (82) supplies both. For  $\bar{x}$  an inertia *lower* bound is wanted, which is exactly what (82) gives by discarding  $J_{CoM} \geq 0$  — and there the mass cancels outright,

$$C_2^2 = \frac{Wz_{CoM}}{J_P} \leq \frac{Wz_{CoM}}{m(z_{CoM}^2 + a^2)} = \frac{g z_{CoM}}{z_{CoM}^2 + a^2}, \quad a \triangleq l_p + p_{\text{off}},$$

so the ceiling on  $C_2$  is pure geometry, with no weight and no inertia in it. It is largest at the low end of  $z_{CoM} \in [0.2, 0.3]$  m (the unconstrained maximum  $\sqrt{g/2a}$  sits at  $z_{CoM} = a$ , below that interval), giving  $C_2 \leq 5.47 \text{ s}^{-1}$  for roll and 5.87 for pitch, hence  $Wz_{CoM} C_2 \leq 47.8$  and 49.7 N m/s at the top of the  $z_{CoM}$  range.

For  $\Delta M_{\text{win}}$  an inertia *upper* bound is wanted, which (82) as written does not give:  $\tau_{\text{end}}$  of (93) grows with  $J_P$ . Retaining the term (82) discards closes it,  $J_P \leq J_{CoM}^{\text{CAD}} + m(z_{CoM,\text{max}}^2 + a^2)$ , i.e. 0.426 and 0.398 kg m<sup>2</sup>. A second and entirely inertia-free ceiling is available as a cross-check, since the excitation window is truncated at the moment cap:  $\Delta M_{\text{win}} \leq M_{\text{cap}} - |M_{\text{crit}}|_{\text{min}}$ , which is 1.67 and 2.34 N m. The first is the tighter and is the one used.

Two further points fix the evaluation. First, the admissible weight is a range, not a number:  $W \in [30.08, 31.59]$  N from unloaded to fully ballasted, and each quantity is read at whichever end is unfavourable for it —  $W_{\text{max}}$  for  $Wz_{CoM}$ ,  $\rho_{\varphi,\text{max}}$  and  $J_P$ ,  $W_{\text{min}}$  for  $\omega_{\text{nom}}$  and for converting a threshold error into an offset error. Stating the box at the unloaded weight alone would leave the ballasted configurations outside it, since  $Wz_{CoM}$  then reaches 9.48 rather than 9.02 N m. Second, at the tilt cap  $x$  is *determined* by the geometry and the ramp rate through  $\Lambda(x) = \varphi_{\text{max}} Wz_{CoM} C_2 / \dot{M}$ , and so is  $\Delta M_{\text{win}} = \dot{M}x / C_2$ ; (92)–(93) exist only to avoid solving that transcendental equation, and must not be put into the denominator of the relative form, which would flatter it. The table below solves for  $x$  once and uses the window it implies, quoting  $\bar{x}$  alongside to show what the shortcut costs.

With  $\varphi_{\text{max}} = 10^\circ$ ,  $\beta_M = 0.0345 \text{ rad}^{-1}$  and  $\dot{M} \in [0.10, 1.20]$  N m/s, the whole of (90)–(97) then closes on the box of Sec. VI-D:

| $\dot{M}$ [N m/s]                            | roll ( $M_x$ ) |       |       | pitch ( $M_y$ ) |       |       |
|----------------------------------------------|----------------|-------|-------|-----------------|-------|-------|
|                                              | 0.10           | 0.45  | 1.20  | 0.10            | 0.45  | 1.20  |
| $x$ at the cap                               | 5.18           | 3.80  | 2.99  | 5.21            | 3.83  | 3.02  |
| $\bar{x}$ by (92)                            | 7.94           | 4.81  | 3.47  | 8.04            | 4.87  | 3.51  |
| $\Delta M_{\text{win}}$ [N m]                | 0.103          | 0.339 | 0.712 | 0.099           | 0.329 | 0.692 |
| $\rho_{\varphi,\text{max}}$ by (101) [mN m]  |                | 68.40 |       |                 | 54.01 |       |
| $\rho_{GE,\text{max}}$ [mN m]                | 0.62           | 2.04  | 4.29  | 0.60            | 1.98  | 4.16  |
| $\bar{\rho}$ by (95) [mN m]                  | 6.99           | 8.64  | 9.92  | 5.66            | 7.04  | 8.15  |
| envelope only, by (90) [%]                   | 396            | 93    | 38    | 335             | 79    | 32    |
| $ e_\omega /\omega_{\text{nom}}$ by (96) [%] | 35.7           | 10.1  | 4.6   | 30.0            | 8.6   | 3.9   |
| $ e_\omega $ [deg/s]                         | 19.8           | 6.2   | 3.2   | 17.3            | 5.4   | 2.8   |
| $ \Delta M_{\text{crit}} $ by (97) [mN m]    | 10.3           | 10.6  | 11.0  | 8.2             | 8.5   | 9.0   |
| $ \Delta \lambda_{\text{off}} $ [mm]         | 0.34           | 0.35  | 0.37  | 0.27            | 0.28  | 0.30  |

The envelope row is the one to read first: at the slowest ramp it exceeds the signal it bounds by four-fold, and only past 0.45 N m/s does it become informative at all. The Chebyshev row below it is a factor of 8 to 11 smaller, and the threshold row is nearly flat in  $\dot{M}$  as the  $x$ -free form requires, only the weak  $\rho_{GE}$  dependence surviving. Its worst value over the box is 11.0 mN m for roll and 9.0 for pitch, i.e. 0.37 and 0.30 mm of identified offset, of which the gravity channel supplies 92%. Here  $\rho_{\varphi,\text{max}}$  is the exact second-order remainder of the gravity moment rather than its Taylor bound,

$$\rho_\varphi = G(\varphi) - G(0) - G'(0)\varphi = W[a(1 - \cos \varphi) + z_{CoM}(\sin \varphi - \varphi)] \leq \frac{1}{2}Wa\varphi_{\text{max}}^2, \quad (101)$$

the second term being negative; the customary bound on the right is loose by 8–14% over the box and there is no reason to pay it.

**The channel the reviewer asks about is the one best absorbed.** One asymmetry between the two channels is worth recording, because it runs opposite to the concern that motivates them. The estimator span contains  $\delta\varphi$ , whose nominal shape is  $\Lambda$ , and the bilinear ground-effect channel goes as  $s\Lambda(s)$  — close to a multiple of a span member, so a single scaling already accounts for all but 10% of it and the full span for all but 4%. The gravity remainder goes as  $\Lambda(s)^2$ , whose exponent is twice as large, and the span cannot follow it: 18% survives. The ground-effect contribution to (97) is thus small twice over — once because  $\rho_{GE,\text{max}}$  is itself an order below  $\rho_{\varphi,\text{max}}$ , and again because what remains of it after the fit re-tunes its own coefficients is smaller still.

Two remarks on robustness. First,  $\bar{x}$  depends on the geometry only as  $(Wz_{CoM}C_2/\dot{M})^{1/3}$ , so a factor-of-two error in  $Wz_{CoM}$  moves it by 26% — the chain is least sensitive to exactly the quantity that is least certain. Second, the  $C_2$  ceiling above is tight enough to be worth checking against the fits, and one of the ten, 5.667 s<sup>−1</sup> for **case\_05**/ $M_x$ , exceeds the roll ceiling 5.468 by 3.6%. The margin is set by  $p_{\text{off}}$ : taking the arm of (82) as  $l_p$  alone raises the ceiling to 5.737 and restores the inequality, which is also the physically correct reading, since  $p_{\text{off}}$  bounds the offset rather than displacing the contact point. The ceiling above is evaluated with  $l_p + p_{\text{off}}$  only because doing so lowers it and is therefore conservative for  $\bar{x}$ .

**Evaluation at the flown configurations.** The table above is the guarantee over the admissible box. The runs themselves sit well inside it, and reading the same bound at the excursions each window actually reached gives the sharper picture below. Here every quantity must come off the *same* window:  $\rho_{\varphi,\max}$  and  $\rho_{GE,\max}$  at the excursion that window reaches,  $\Delta M_{\text{win}}$  as its own moment increment. Over the 140 gate-passing runs, with the loaded weight  $W = 31.59$  N,  $l_p + \lambda_{\text{off}} = 0.160/0.130$  m (roll/pitch),  $z_{CoM} = 0.30$  m and  $\beta_M = 0.0345$  rad<sup>-1</sup>:

| $\dot{M}$ [N m/s]                       | 0.10  | 0.30  | 0.65  | 1.20  | median | worst |
|-----------------------------------------|-------|-------|-------|-------|--------|-------|
| $ \varphi _{\text{end}}$ [deg]          | 5.23  | 5.90  | 6.55  | 6.17  | 5.95   | 9.33  |
| $\Delta M_{\text{win}}$ [N m]           | 0.072 | 0.165 | 0.306 | 0.479 | 0.221  | —     |
| $\rho_{\varphi,\max}$ [mN m]            | 19.7  | 22.5  | 30.4  | 27.2  | 23.6   | 58.3  |
| $\rho_{GE,\max}$ [mN m]                 | 0.24  | 0.59  | 1.24  | 1.85  | 0.68   | 3.00  |
| $x = C_2 \tau_{\text{end}}$             | 3.47  | 2.64  | 2.15  | 1.79  | 2.67   | 5.20  |
| $ e_\omega /\omega_{\text{nom}}$ [%]    | 13.8  | 6.0   | 4.7   | 2.9   | 4.9    | 25.1  |
| $ \Delta M_{\text{crit}} $ [mN m]       | 1.94  | 2.73  | 3.77  | 3.76  | 3.01   | 7.86  |
| $ \Delta \lambda_{\text{off}} $ [mm]    | 0.061 | 0.086 | 0.119 | 0.119 | 0.095  | 0.249 |
| ground-effect share of $\bar{\rho}$ [%] | 1.8   | 4.0   | 6.0   | 8.8   | 5.1    | 25.6  |

These too are bound evaluations, not measurements: each row forwards the *unprojected* channel maximum through (96). Using instead the  $x$ -free constants of (97) the same runs give  $|\Delta M_{\text{crit}}| \leq 8.8$  mN m, i.e.  $|\Delta \lambda_{\text{off}}| \leq 0.28$  mm — the price of not evaluating  $R_\varphi, R_{GE}$  is under 12%, and the flown excursions buy a further 27% over the 12.0 mN m the box guarantees.

The per-run propagation reported with the error budget is a separate and much smaller number, because it convolves the *measured*  $\rho$  with the exact kernel after the estimator has absorbed its own span. Over the same runs the residual forcing that actually reaches the onset is 2.75 mN m in the median (12.28 at worst) on the gravity channel and 0.53 (3.20) on the ground-effect channel, the rate deviation  $|e_\omega|/\omega_{\text{nom}}$  is 0.24% in the median (6.5% at worst summed over channels), and the resulting threshold error is 0.04 mN m in the median and 0.15 at worst — 0.001 to 0.004 mm of offset. The two differ by roughly a factor of 50, which is the price the envelope form pays for assuming nothing about  $\rho$  but its size.

The bilinear ground-effect channel contributes 5.1% of  $\bar{\rho}$  in the median and 25.6% at worst; it grows with the ramp rate because  $\rho_{GE} \propto \dot{M}\tau$ , whereas the gravity remainder depends on the excursion alone. Its own contribution to the threshold is 0.00 mN m in the median and 0.01 at worst once propagated, so the channel the reviewer asks about is not the one that sets the budget. **The threshold bound  $|\Delta \lambda_{\text{off}}| \leq 0.25$  mm is the figure that matters: it does not involve  $\Delta M_{\text{win}}$ , and so does not inherit the loss the relative normalisation suffers at slow ramps.** The relative bound is loosest at 0.10 N m/s, where the window closes on a small moment increment while the excursion is unchanged — the same ordering the excitation design reports — and it is conservative by roughly a factor 5 overall, since it forwards the *unprojected* remainder (23.6–58.3 mN m here, against 1.2–12.3 mN m measured after the estimator absorbs its own degrees of freedom).

## VI-B: ramp input tracking (rate limits)

The rotor response identified in Sec. V (Table 4:  $\omega_{n,\text{rot}} = 21.74$  rad/s,  $\zeta = 0.7814$ ,  $\tau_d = 15.67$  ms) tracks a ramp with the steady-state lag

$$t_{\text{lag}} = \frac{2\zeta}{\omega_{n,\text{rot}}} + \tau_d = 71.9 + 15.7 = 87.6 \text{ ms},$$

so the realised moment trails the command by  $\dot{M} t_{\text{lag}}$ , i.e. 8.8 mN m at 0.10 N m/s and 105 mN m at 1.20 N m/s. This lag does *not* bias  $M_{\text{crit}}$ . The identification consumes the *realised* moment, reconstructed from the measured rotor speeds, not the commanded one; a common lag therefore shifts the moment trace and the rate trace together and leaves the moment read at the detected onset unchanged. What the lag does consume is window length, which the sample-count gate enforces.

Two residual effects remain and are both gated. First, the realised *slope* may differ from the commanded rate; the run-level gate rejects  $|\dot{M}_{\text{meas}} - \dot{M}_{\text{cmd}}|/\dot{M}_{\text{cmd}} > 3\%$ , and since  $C_1 = K\dot{M}$ , a 3% rate error perturbs the amplitude constant by 3% — well inside the  $\pm 20\%$  mis-specification the sensitivity analysis shows to be harmless. Second, the start of the ramp is not linear while the rotor transient decays; with  $4/(\zeta\omega_{n,\text{rot}}) = 235$  ms to 2%, that transient is confined to the low-moment part of the window, which the slope gate excludes

by scoring only the segment above the per-axis moment floor. The linearity gate (30 mN m RMSE about the best-fit line) then acts as a fault detector for ramps that are stepped or aborted outright.