

Sec. VIII-x: the error analysis, verified on the campaign

drop-in; figure `fig.sec8_error.png`

0.1 What is guaranteed, and what is measured

Two different kinds of statement are made about the accuracy of the identification and they must not be merged. Equation (108) is a *guarantee*: given a scale, a CAD inertia and the excitation protocol, the terms the model neglects cannot displace the identified threshold by more than $\bar{\rho}$, which is 12.04 mN m for roll and 9.95 for pitch, or 0.400 and 0.322 mm of equivalent offset. It says nothing about sensor noise, rig repeatability or anything else it did not model. The repeatability reported below is a *measurement* of the whole error, with no model in it at all. Quoting either as if it were the other would misstate both.

0.2 The bound, checked forwards

The guarantee is verified as described in Sec. ??: each perturbation is evaluated along the measured trajectory, propagated through the deviation dynamics and correlated against the weight of (107). All 140 runs fall inside, with a median displacement of 0.049 mN m, 0.161 at p90 and 0.927 at worst against the 12.04 allowed.

That margin is not an accident and it has a closed form. Writing $x = C_2\tau_{\text{end}}$, the gravity remainder behaves as $\rho_\varphi \propto \varphi^2 \sim e^{2C_2s}$ near the window end, so $e_\omega \sim e^{2x}$ while $\bar{\rho} \sim e^{2x}/x$ and therefore the envelope of (93) grows as $E \sim e^{3x}/x$:

$$\frac{|e_\omega|}{E} \sim x e^{-x}. \quad (\text{VIII.1})$$

Both grow with the window; the bound grows by one whole factor of e^x more, because the Chebyshev step of (93) puts the window average of ρ underneath a kernel that is largest exactly where ρ is smallest. Equivalently, keeping the covariance that Chebyshev discards turns the inequality into an identity,

$$\frac{|e_\omega|}{E} = 1 + \text{corr}(\rho, h_\omega) \text{CV}_\rho \text{CV}_h, \quad (\text{VIII.2})$$

and lengthening the window makes both coefficients of variation grow faster than the correlation weakens. Measured, the realised deviation is 3.2% of the bound at $\dot{M} = 0.10$ and 20.8% at 1.20; the reduction factors evaluated at their suprema account for a further factor of 1.5 to 2.1, roughly constant across the rates.

This says how informative the guarantee is and nothing else, and the distinction is worth making explicitly because the two are easily conflated. A larger occupancy means the bound is closer to the truth, not that the truth is larger or smaller. On the ρ channel the fast ramp is in fact slightly *worse*: the deviation is less aligned with the onset direction there (31.4° against 23.0°, so 82% of it is absorbable rather than 89%), but $\Delta M_{\text{crit}} = \dot{M} \delta$ and δ falls only 2.3× across the rate range while $\dot{M}$ rises twelvefold, so the channel contributes 0.35 mN m at $\dot{M} = 0.10$ and 1.82 at 1.20. The recommendation on excitation rate is made separately below, on measured grounds, and does not rest on this paragraph.

0.3 The fit against its cap

Because the nominal response is itself an admissible member of the fitted family, the minimiser cannot do worse than it does, and with the sensor noise added by the triangle inequality

$$\|\hat{r}\| = \min_g \|y - g\| \leq \|y - f\| \leq \|e_\omega\| + \|n\| \leq \|E\| + \|n\|. \quad (\text{VIII.3})$$

The comparison needs the residual split first. e_ω is a Duhamel integral of a smooth ρ and carries essentially nothing above a few times $1/\tau_{\text{end}}$, so residual content faster than that cannot be the small-angle or ground-effect remainder whatever else it is; charging the model for it is not a test of the model. Splitting at 5 Hz (the windows are 53 samples at 100 Hz, too short for an IIR filter), the low-frequency residual lies inside the

Sec. VIII: the error of the cosh method, verified on the campaign




Figure 1: The error of the identification, measured on the 140-run campaign. (a) the identified critical moment against the ramp rate for all twenty configuration/direction groups, each scaled by its own mean; M_{crit} is a static property of the rig, so the ideal is a flat line and every departure is error, with no model in the comparison. (b) the same pooled by rate: only the slowest ramp departs significantly. (c) the domain test — σ_M is flat and $\sigma_t = \sigma_M/\dot{M}$ is not, so the dominant error is in the moment rather than in the located onset. (d) the fit residual against its a priori cap, split at 5 Hz; only the low-frequency part can be the modelling remainder. (e) the reported half-sum, repeated at seven rates per configuration, against the bound of (108). (f) the budget, bounded terms on the left and measured ones on the right.

cap for 114 of 140 runs, the 26 exceptions all at the two fastest rates and by no more than a factor 1.16. Including the high-frequency content the figure is 88/140.

The split also identifies what the rest is. Above 5 Hz the residual carries $1.02^\circ/\text{s}$ while the pre-onset record carries 0.14, a factor of seven: before the onset the vehicle is supported by every landing gear and after it pivots on one edge with far less damping, so the pre-onset floor understates the in-window disturbance and must not be used as $\|n\|$.

0.4 Repeatability, with no model in it

The campaign supplies a model-free accuracy measurement, because M_{crit} is a *static* property of the rig — the moment at which the vehicle is on the verge of tipping over a landing-gear edge — and cannot depend on how fast the moment was ramped to reach it. Identifying it at seven rates per configuration and per direction therefore repeats a measurement whose true value is known to be constant, and Fig. 1(a) is that test.

	per direction	reported half-sum
scatter σ_M	21.3 mN m	—
as equivalent offset	0.71 mm	—
1 σ across the rates	—	0.28 to 1.13 mm
median	—	0.522 mm
median, $\dot{M} \geq 0.65$	—	0.327 mm

The two columns are consistent with each other without adjustment: two directions with independent moment-domain errors give a half-sum spread of $\sigma_M/\sqrt{2} = 0.50$ mm, and the measured median is 0.522.

Which domain. A mislocated onset is an error in *time*, so it would show a flat σ_t and a σ_M rising with $\dot{M}$; a vehicle that does not tip at exactly the same moment twice is an error in *moment*, so σ_M is flat and

σ_t falls as $1/\dot{M}$. Measured, σ_M runs 13.7 to 33.5 mN m with a coefficient of variation of 34% while σ_t runs 11.4 to 225 ms at 97%: the error is in the moment.

The alternative can be excluded quantitatively rather than by inspection. If the scatter were a within-run disturbance seen through the onset, then for a disturbance of fixed amplitude $\sigma(\Delta M_{\text{crit}}) \propto \sqrt{C_2/B(x)}$, so the short window — the fast ramp — would scatter *more*, by a factor 5.72 across the rate range. It scatters 0.70, an eightfold disagreement in the opposite direction. The scatter is a run-to-run change in where the vehicle actually tips: landing-gear settling, floor friction and repositioning between runs.

The slowest ramp. One condition departs systematically. Pooling the identified M_{crit} less each configuration's own mean, $\dot{M} = 0.10$ N m/s reads -26.70 mN m against $+4.45$ for the rest (Welch $t = -4.34$, $p = 0.0002$), with 75% of those runs below their own mean. As an onset error that is 267 ms, matching the 225 ms of σ_t measured there, and its sign says the onset is found *early* — what a slow rise emerging from noise does to a residual minimiser. It is not ρ , which contributes 0.3 to 1.8 mN m on these channels, twenty times too small and with the opposite rate trend. Restricting the protocol to $\dot{M} \geq 0.65$ N m/s removes it and takes the reported spread from 0.522 to 0.327 mm.

0.5 What the two-sided difference removes

The offset is read from the half-sum $M_{\text{off}} = (M_+ + M_-)/2$, and the two runs tip over different edges with restoring arms $A_{\pm} = l_p \mp a$. In mirrored coordinates both obey the same equation, so in signed φ the arm term of the deviation forcing carries $\text{sign}(\dot{M})$ while the $Wz \sin \varphi$ term does not: the arm term is antisymmetric and cancels. Simulating the full nonlinear tip-over with a 2 mm offset and fitting with the deployed estimator, 98.70% of the per-direction bias cancels at every rate, leaving about 1 μm . The ground-effect channel cancels as well: β_M flips sign with the pivot edge — it is proportional to $\sum_i (l_{y,i} + l_p)^2 b_{i,\varphi}^\dagger$, and the allocation basis reverses — so the bilinear term of (79) is odd and contributes 0.6 to 0.9 μm instead of the 46 to 78 it would if it were even. The per-configuration calibration cancels too, being common to both directions by construction: $\pm 10\%$ of amplitude error and $\pm 5\%$ of exponent error are rejected to 99.77–99.98% and move the offset by 0.000 mm.

What does not cancel is the weight, because (34) reads the *difference* of the two critical moments, where the same bias adds. Per 1% of calibration error that is 0.027% of W from the amplitude channel and 0.077% from the exponent, the exponent being the more damaging by about three; the closed forms are $A(x)/B(x)$ and $D(x)/B(x)$ times $\dot{M}/C_2$.

0.6 Choosing the excitation rate

The protocol's seven rates permit the choice to be made on measurement rather than argument. Taking each configuration's half-sum at a given rate against its own mean over all rates isolates the error contributed by that rate:

$\dot{M}$ N m/s	post-onset samples	window s	RMS error mm	p90 mm
0.10	74	0.740	0.810	1.035
0.20	64	0.635	0.569	0.975
0.30	58	0.575	0.708	0.994
0.45	54	0.545	0.707	1.374
0.65	48	0.480	0.528	0.924
0.90	44	0.445	0.443	0.696
1.20	41	0.410	0.422	0.637

There is no interior optimum: the error falls to the fastest rate tested. That is worth checking against the obvious objection, which is that a faster ramp shortens the window and leaves fewer samples to fit. It does — 74 down to 41 — but the onset is located by $\|\chi\|^2 = C_1^2 C_2 B(x)$ with $C_1 = \dot{M}/Wz_{CoM}$, so the amplitude rises twelvefold while the sample count falls only $1.8\times$, and the product improves: $\|\chi\|$ grows $1.35\times$ and the onset uncertainty σ_δ falls by 26%. Forty-one samples remain ample for a two-parameter fit.

Two effects would eventually create an interior optimum and neither has taken over by 1.20 N m/s. The onset is quantised to the sample grid, contributing $T_s \dot{M}/\sqrt{12}$ per direction, or 0.081 mm to the half-sum at the fastest rate against the 0.422 measured — under 4% of the variance, and removable outright by refining the

onset below the grid. And the pseudo-inverse allocation saturates at $M_{x,\max} = 2.371$ and $M_{y,\max} = 2.737$ N m (Figs. 13–14), which caps the usable rate from above. Within the range tested the recommendation is simply the top of it, and restricting to $\dot{M} \geq 0.65$ N m/s takes the reported spread from 0.522 to 0.327 mm.

0.7 Summary

term	how known	equivalent offset
small-angle + GE remainder	bounded a priori, (108)	≤ 0.400 mm
realised, p90	forward check, 140/140	0.005 mm
onset location	not resolved; below the next line	—
rig repeatability, per direction	measured	0.71 mm
reported offset, all rates	measured, $7\times$ per config	0.522 mm
reported offset, $\dot{M} \geq 0.65$	measured	0.327 mm

The modelling approximations are therefore a bounded minority of an error the rig dominates: whatever limits this method, it is not the neglected $\sin \varphi$. That is worth stating plainly because it also justifies the change of onset model. The quadratic growth law used previously carried a *systematic* error of 12 to 31% of the window in the located onset — 22 to 58 mN m over the observed x range, comparable to or larger than the entire error measured here, and a bias rather than a scatter, so averaging does not reduce it. Retaining $Wz_{CoM}\varphi$, which is exactly what turns the quadratic into the hyperbolic cosine with $C_2^2 = Wz_{CoM}/J_P$, moves the algorithm’s contribution from dominant to negligible.