

COSH: Closed-Form Onset Search via Hyperbolic Response

Methodology for pre-flight critical-moment identification of a multirotor from ground-contact tip-over excitation

This document specifies the complete identification methodology, from raw signals to the take-off feed-forward moment. Every stage mirrors the implementation in `critical_value_getter_pieewise.py`; the companion analyses (model-fidelity bounds, excitation design, ground-effect bracket) are documented in `fidelity_section.tex` and are only referenced here.

1 Problem setting and notation

A multirotor rests on its landing gear. Ramping the commanded roll (pitch) moment tilts the vehicle until it begins to rotate about the landing-gear contact line — the *pivot*. The applied moment at that *rotation onset* is the critical moment M_{crit} ; comparing the two tip directions yields the CoM offset and the feed-forward moment M_{ff} that compensates take-off. No rig, no motion capture and no stable in-flight data are used.

Symbol	Meaning	Value	Provenance
W, z_{CoM}	weight, CoM height	31.59 N, ~ 0.3 m	measured / conservative
$l_{p,\phi}, l_{p,\theta}$	pivot offsets (roll, pitch)	0.140, 0.110 m	conservative geometry
J_P	inertia about the pivot line	—	never used numerically
C_T	thrust map, $T_i = C_T \Omega_i^2$	1.3175×10^{-7}	bench calibration
M_{thr}	window threshold	0.01 N m	noise floor
M_{cap}	allocator caps (x, y)	2.37, 2.74 N m	allocation feasibility
ϵ_{max}	ramp-slope gate	3%	sensitivity chain (Sec. 4)
N_{min}	window-sample gate	38	30 pre + 8 post
$\epsilon_{\text{lin,max}}$	linearity-RMSE gate	30 mN m	$10 \times$ noise floor
N_{seg}	sweep margin (each end)	8	degenerate-candidate guard

Table 1: Design constants. (C_2, K) are *calibrated*, not set (Sec. 6).

2 Contact-phase dynamics and the closed form

About the active pivot line, with the small-angle attitude dependence retained ($\cos \varphi \approx 1$, $\sin \varphi \approx \varphi$), the roll/pitch dynamics read

$$J_P \dot{\omega} = M(t) + f l_p - W (l_p + \lambda_{\text{off}}) + W z_{\text{CoM}} \varphi, \quad (1)$$

where M is the commanded body moment, f the collective thrust and λ_{off} the CoM offset along the excited axis. Below the critical value the vehicle is static. At the onset the net moment vanishes, which supplies *two* anchoring conditions,

$$\omega(t_{\text{crit}}) = 0, \quad \dot{\omega}(t_{\text{crit}}) = 0. \quad (2)$$

Past the onset gravity feeds back positively: the linearized dynamics have *real* eigenvalues $\pm C_2$ and, with the measured moment ramp $M(t) = M_{\text{crit}} + \dot{M} \tau$ in onset-relative time $\tau = t - t_{\text{crit}}$, the exact solution of (1)–(2) collapses to the hyperbolic closed form

$$\omega(\tau) = C_1 (\cosh(C_2 \tau) - 1), \quad C_1 = K \dot{M}, \quad K = \frac{1}{W z_{\text{CoM}}}, \quad C_2 = \sqrt{\frac{W z_{\text{CoM}}}{J_P}} \quad (3)$$

for $\tau \geq 0$ (and $\omega = 0$ for $\tau < 0$). Because $J_P C_2^2 = W z_{\text{CoM}}$, the amplitude C_1 is *fixed* by the measured ramp rate: with (C_2, K) pinned as rig constants and the baseline fixed by continuity, **no shape parameter is free per run**. The time-quadratic model of the original submission is the $C_2 \rightarrow 0$ limit of (3) — it discards the gravity term that creates the instability, with 14–75% truncation error over the windows actually used — and is retained only as a baseline.

3 Signals and excitation window

Per run, the moment is reconstructed from the measured rotor speeds, $M(t) = C_T \sum_i b_{a,i} \Omega_i^2(t)$, and $\omega(t)$ is taken from the odometry rate. The window end is $i_{\text{end}} = \arg \max_t |M|$, *truncated at the first sample exceeding the allocator cap* M_{cap} : past it the thrust allocation saturates, the executed ramp is no longer linear, and the cosh family does not apply. (On the reference dataset the cap never binds; peak $|M| \leq 1.7$ N m.) The window start is the *last* sub-threshold sample before the peak, walking backwards — robust to spin-up blips that cross the threshold long before the actual ramp.

4 Onset-free run gates

Before any onset is estimated, each run must pass three gates evaluated on the measured moment trace alone (hence free of circularity):

1. **Ramp-slope error** $|\varepsilon| \leq 3\%$, with $\varepsilon = (|\dot{M}| - \dot{M}_{\text{cmd}})/\dot{M}_{\text{cmd}}$ and $\dot{M}$ the least-squares slope over the window. A 3% error in $\dot{M}$ perturbs $C_1 = K\dot{M}$ by 3% — about $7\times$ below the $\pm 20\%$ level the sensitivity analysis shows harmless.
2. **Window samples** $N_{\text{full}} \geq 38$: a necessary condition for any onset position. Because the window opens at $|M| = 10$ mN m while $M_{\text{crit}} \gtrsim 0.4$ N m, the onset cannot occur earlier than $\hat{N}_{\text{pre}} = M_{\text{crit}}^{\text{min}}/(\dot{M} T_s) \geq 30$ samples into the window even at the fastest ramp; adding the $N_{\text{seg}} = 8$ post-onset fit minimum gives 38.
3. **Linearity RMSE** ≤ 30 mN m about the run’s own best-fit line: a fault detector for aborted/stepped ramps, set at $10\times$ the execution noise floor and above the healthy-run maximum.

The gates run *before* calibration so a poorly executed ramp cannot contaminate the dataset-coupled constraint $C_1 = K\dot{M}$.

5 Onset detection

With (C_2, K) pinned, detection reduces to an exhaustive sweep of the single remaining unknown — the onset index. Each candidate is a plain arithmetic evaluation: no optimizer, no initial guess, no seed model.

Algorithm 1 COSH onset detection (single run)

Require: samples $\{(t_i, \omega_i)\}$ on the excitation window $[i_{\text{start}}, i_{\text{end}}]$; moment $M(t)$; calibrated (C_2, K) ; measured ramp slope $\dot{M}$; margin $N_{\text{seg}} = 8$

Ensure: onset index j^* ; critical moment M_{crit}

- 1: $C_1 \leftarrow K\dot{M}$ ▷ onset conditions (2) fix the amplitude
 - 2: **for** $j = i_{\text{start}} + N_{\text{seg}}, \dots, i_{\text{end}} - N_{\text{seg}}$ **do** ▷ exhaustive sweep of the single unknown
 - 3: $C \leftarrow \text{median}\{\omega_i : i_{\text{start}} \leq i < j\}$ ▷ baseline by continuity
 - 4: $\text{cost}(j) \leftarrow \sum_{i=i_{\text{start}}}^{j-1} (\omega_i - C)^2 + \sum_{i=j}^{i_{\text{end}}} (\omega_i - [C_1(\cosh C_2(t_i - t_j) - 1) + C])^2$
 - 5: **end for**
 - 6: $j^* \leftarrow \arg \min_j \text{cost}(j)$; $M_{\text{crit}} \leftarrow M(t_{j^*})$
-

Two implementation details. *Baseline by median.* The pre-onset gyro trace is heavy-tailed (propeller vibration, spin-up spikes), and — more importantly — when a candidate j overshoots the true onset, the pre-segment tail contains the beginning of the hyperbolic rise. A mean baseline drifts with that leaked rise and partially cancels the residual it should expose, flattening the cost landscape; the median stays at the rest level until half the segment is contaminated, so misplaced candidates are penalized sharply and

the minimum stays well defined. It is also a one-shot, seedless statistic, consistent with the no-optimizer design. An ablation rerunning the full pipeline — including the (C_2, K) calibration — with the mean instead (`analysis/baseline_stat_ablation.py`) leaves the run gating unchanged and 62/132 onsets identical, but biases the remainder toward *earlier* onsets, shrinking the directional means in magnitude by up to 30 mNm on the noisiest datasets. The route is indirect. At *fixed* (C_2, K) a baseline biased toward the motion direction *delays* the detected onset: the post-onset residual $C_1[\cosh(C_2(t - t_{\text{crit}}^{\text{real}})) - \cosh(C_2(t - t_{\text{crit}}^{\text{est}}))] - C$ can only vanish with a later split, and an oppositely biased baseline advances it. The dominant effect, however, is that the leaked-rise flattening of the cost landscape moves the CV-selected calibration to a different point on the (C_2, K) ridge — e.g. $C_2: 8.00 \rightarrow 5.38$, $K: 0.060 \rightarrow 0.150$ on the most affected dataset — whose net effect is the earlier-onset magnitude shrink; the ramp-invariance CV itself is essentially unchanged. The pivot-free offset M_{ff} moves by at most 4.2 mNm (0.13 mm), well below the run-to-run scatter — the median removes a systematic bias and tunes nothing. A second ablation pins $C = 0$, the ideal-sensor value implied by the onset conditions: the measured pre-onset baseline is in fact nonzero (gyro bias plus residual spin-up motion; $|C|$ median 3.1, p90 14.2, max 30.6 mrad/s), yet the estimate barely moves (71/132 onsets identical, M_{ff} shift ≤ 1.3 mNm). The estimated baseline is therefore not load-bearing on this dataset; it is kept because it makes the sweep invariant to the gyro bias — a rig- and temperature-dependent quantity that need not stay this small — at zero parameter cost. Nor does C need the bias to be constant in time: it is re-estimated per run, so only *within-window* drift (windows last 1–10 s) matters. Measured over the guaranteed pre-onset segment ($|M| < 0.3$ N m, spans ≤ 3.0 s), the baseline wanders by 3.4 mrad/s (median; p90 24, max 68) — an upper bound that conflates true sensor drift with real elastic pre-onset motion (`analysis/bias_drift.py`). This wander is present in the very data on which the $C = 0$ bound and the rate-invariance certificates were earned, and its structure is covered by the perturbation absorption: the constant part goes to C , a τ -linear trend is absorbed by the span projection, and only the residual curvature reaches ρ . *Sweep margin*. The margin $N_{\text{seg}} = 8$, applied at both window ends, guarantees every candidate split leaves at least N_{seg} samples in each segment; it excises only degenerate candidates: near $j = i_{\text{start}}$ the baseline would be estimated from a handful of samples, and near $j = i_{\text{end}}$ the post-onset arc is too short for the cosh branch to be falsifiable (at $j = i_{\text{end}}$ its cost is nearly an empty sum). Since physically feasible onsets lie at least $\hat{N}_{\text{pre}} \geq 30$ samples inside the window (Sec. 4) and the window closes only after the response is well developed, the sweep remains exhaustive over every feasible onset.

Remark 1 *Because the model curve is fully determined before the post-onset data are seen, goodness of fit is a prediction test, not flexibility: across all gate-passing runs the normalized post-onset residual is $\approx 6\%$ (gyro-noise dominated) and invariant over the twelvefold ramp-rate range — the empirical certificate that the response stays in the cosh family.*

6 Rig-constant calibration

Algorithm 1 presupposes (C_2, K) . Although both have geometric expressions (3), they are *not* computed from geometry: z_{CoM} and especially J_P are not precisely known a priori, and — more importantly — the constants that make (3) reproduce the measured response are *effective* values that absorb the state-linear contributions of the unmodeled effects (the local slope of the gravity nonlinearity and the ground-effect spring). Both are calibrated once per dataset.

Joint residual minimization is not a usable criterion, and the reason can be made exact. The identity $\cosh x - 1 = 2 \sinh^2(x/2)$ factors the closed form as

$$C_1(\cosh(C_2\tau) - 1) = \underbrace{\frac{C_1}{2} e^{-C_2 t_{\text{on}}}}_A e^{C_2 t} (1 - e^{-C_2 \tau})^2, \quad (4)$$

whose last factor rises to one within a few growth times C_2^{-1} . On the tail the amplitude and the onset therefore enter the data only through the product A , and the substitution

$$C_1 \rightarrow \lambda C_1, \quad t_{\text{on}} \rightarrow t_{\text{on}} + C_2^{-1} \ln \lambda \quad (5)$$

leaves A *exactly* invariant for every $\lambda > 0$: the response changes only through the last factor, by the relative amount $[(1 - \lambda e^{-C_2 \tau}) / (1 - e^{-C_2 \tau})]^2 - 1 \approx -2(\lambda - 1) e^{-C_2 \tau}$, which decays exponentially in τ . An earlier onset with a smaller amplitude thus trades off against a later onset with a larger one everywhere except near the onset — where the response is quadratic, $\frac{1}{2} C_1 C_2^2 \tau^2$, and of the order of

the gyro noise. Equivalently, the sensitivity vectors $\partial\omega/\partial C_1 \propto \cosh(C_2\tau) - 1$ and $\partial\omega/\partial t_{\text{on}} \propto \sinh(C_2\tau)$ become parallel on the tail, so the Gauss–Newton Hessian of the summed residual is nearly singular along $\delta C_1/C_1 = C_2 \delta t$ — a *shallow ridge* in the (C_2, K) plane whose minimizer is set by noise, not physics. This is an ill-conditioning statement, not a non-identifiability one: in the noiseless limit the post-onset shape determines all three parameters uniquely. Under gyro noise, however, the information along (5) is exponentially small, so per-run estimates — and with them the detected critical moments — scatter along the ridge. The calibration should therefore be read as *variance suppression*: choosing (C_2, K) once per dataset to minimize the dispersion of M_{crit} across ramp rates removes precisely that scatter, which is where the practical strength of the method lies.

The ridge is invisible to the residual but not to the deliverable. Moving along (5) shifts every detected onset by the same *time* offset $\delta t = C_2^{-1} \ln \lambda$ — with $\lambda = K^*/K$ the same for all runs, since $C_1 = KM$ scales them alike — and therefore biases the critical moment by

$$\delta M_{\text{crit}} = \dot{M} \delta t \propto \dot{M}, \quad (6)$$

linearly in the ramp rate: a wrong pair *fans out* the identified M_{crit} across the twelvefold rate range, while the correct pair leaves it rate-invariant. The direction the residual cannot see is thus exactly the direction a rate-consistency criterion sees best. We therefore calibrate by *physical self-consistency*: a static threshold must not depend on how fast it was approached. With runs at ramp rates $r \in \{0.10, \dots, 1.20\}$ N m/s,

$$\text{score}(C_2, K) = \sum_{\text{dir} \in \{+, -\}} \text{CV}(\{M_{\text{crit}}^{(r)}(C_2, K)\}_r), \quad \text{CV} = \text{std}/|\text{mean}|, \quad (7)$$

where $M_{\text{crit}}^{(r)}$ is the output of Algorithm 1 on the run with commanded rate r .

Search ranges. The admissible box is set so that it covers every physically plausible constant with margin. For $K = 1/(W z_{\text{CoM}})$, the coarse range $K \in [0.10, 0.50]$ (N m)^{−1} corresponds to $z_{\text{CoM}} \in [0.06, 0.32]$ m at the vehicle weight — bracketing the plausible CoM heights — with the refinement allowed to move within the wider clip box $[0.05, 0.70]$ (z_{CoM} up to 0.63 m). For $C_2 \in [3, 8]$ rad/s, the implied pivot inertia $J_P = 1/(KC_2^2)$ spans roughly $[0.03, 1.1]$ kg m², bracketing the parallel-axis estimate $J + mz_{\text{CoM}}^2 \approx 0.2$ – 0.35 kg m² by a factor of several on both sides.

Algorithm 2 Ramp-invariance consistency score

Require: gated runs at rates $r \in \{0.10, \dots, 1.20\}$ N m/s, both tip directions; candidate pair (C_2, K)

Ensure: consistency score s ▷ lower = M_{crit} more rate-invariant

- 1: **for** each gated run r **do**
 - 2: $M_{\text{crit}}^{(r)} \leftarrow$ Algorithm 1 with (C_2, K) pinned ▷ stride-3 sweep
 - 3: **end for**
 - 4: **for** $\text{dir} \in \{+, -\}$ **do**
 - 5: $\text{CV}_{\text{dir}} \leftarrow \text{std}(\{M_{\text{crit}}^{(r)}\}_{r \in \text{dir}}) / |\text{mean}(\{M_{\text{crit}}^{(r)}\}_{r \in \text{dir}})|$
 - 6: **end for**
 - 7: **return** $s \leftarrow \text{CV}_+ + \text{CV}_-$ ▷ Eq. (7): a static threshold must not depend on the approach rate
-

Algorithm 3 Rig-constant calibration (coarse-to-fine)

Require: gated runs; box $C_2 \in [3, 8]$ rad/s, $K \in [0.05, 0.70]$ (N m)^{−1}

Ensure: calibrated pair (C_2^*, K^*)

- 1: **coarse:** evaluate Algorithm 2 on the 11×11 grid $C_2 = 3.0:0.5:8.0$, $K = 0.10:0.04:0.50$
 - 2: $(C_2^b, K^b) \leftarrow$ arg min over the coarse grid
 - 3: **fine:** evaluate Algorithm 2 on a 9×9 uniform grid over $[C_2^b \pm 0.5] \times [K^b \pm 0.04]$, clipped to the box ▷ resolution 0.125 rad/s, 0.01 (N m)^{−1}
 - 4: $(C_2^*, K^*) \leftarrow$ arg min over the fine grid
-

Coarse-to-fine procedure. Algorithm 2 states the consistency criterion as a stand-alone routine — it wraps Algorithm 1 and returns the directional sum of coefficients of variation — and Algorithm 3 performs the search over it. The coarse scan evaluates (7) on an 11×11 grid ($\Delta C_2 = 0.5$ rad/s, $\Delta K = 0.04$ (N m)^{−1};

121 evaluations). A single refinement pass then re-scans a 9×9 uniform grid spanning $\pm$ one coarse step about the coarse minimizer (± 0.5 rad/s, ± 0.04 (Nm) $^{-1}$, clipped to the box), i.e. final resolutions of 0.125 rad/s and 0.01 (Nm) $^{-1}$ — so the reported constants are not quantized to the coarse step. One pass suffices because the objective is flat along the (C_2, K) ridge; it is also *piecewise constant* (the onset is an integer index, so small parameter changes leave M_{crit} unchanged), which is why gradient-based or simplex optimizers stall on it, and a population-based global search (differential evolution) was found to cost about $4\times$ more for no measurable gain. During calibration the onset sweep of Algorithm 1 is subsampled by a stride of 3 samples for speed; the final identification runs at full resolution with the selected pair. Each evaluation is an arithmetic sweep, so no iterative optimizer is involved anywhere in the pipeline.

Remark 2 *The objective (7) retains a shallow valley, so the individual constants are only loosely determined: $Wz_{\text{CoM}} = 1/K$ and $J_P = 1/(KC_2^2)$ are order-of-magnitude sanity checks, not precision measurements. What the criterion pins down — by construction — is the identified onset: M_{crit} is robust along the valley, and a $\pm 20\%$ mis-specification of either constant moves the identified CoM by less than 0.9 mm.*

7 Aggregation and deliverables

Directional means over the gated runs give $\bar{M}_+$ and $\bar{M}_-$. The take-off feed-forward moment is the *pivot-free average*

$$M_{\text{ff}} = \frac{1}{2} (\bar{M}_+ + \bar{M}_-), \quad (8)$$

in which the pivot-arm terms $\pm(W - f)l_p$ — and any thrust-proportional contamination acting through the arm l_p , ground effect included — cancel by antisymmetry, while landing-gear asymmetry is deliberately *retained*: the same contact-phase moment acts at lift-off, where M_{ff} compensates it. The CoM offset follows from the moment-to-offset conversion of M_{ff} and is used for validation against load-cell ground truth only; M_{ff} itself is the quantity applied in flight. A roll excitation (M_x) senses the y -offset via $M_{\text{ff},x} = +Wy_{\text{off}}$; a pitch excitation (M_y) senses the x -offset via $M_{\text{ff},y} = -Wx_{\text{off}}$.

8 Comparison with per-run nonlinear least squares

The predecessor of Algorithm 1 fitted each run individually by nonlinear least squares (SCIPY `least_squares`, trust-region reflective): free (C_1, C_2, C) with C_2 bounded to the physical band $[3, 8]$ rad/s and a free baseline, the onset swept locally around a quadratic-model seed. To quantify what the dataset-coupled calibration buys, both estimators were run on *identical* inputs — same signals, excitation windows, allocator caps and onset-free gates (the gated run set is method-independent because the gates use the moment trace only), 132 runs in total (`analysis/nls_comparison.py`). Table 2 compares the deliverables — the pivot-free offset (8) and the CoM offset — against the load-cell ground truth; Table 3 compares the directional critical-moment means and their ramp-rate dispersion.

Three observations. (i) *Run-level robustness is where the methods separate*: the directional CV is lower under COSH in 20/20 case-axis-direction groups — median 2.4% vs. 5.6%, NLS worst case 15.2% — and the paired per-run critical moments differ by 27.6 mNm at the median but 118 mNm at the 90th percentile and up to 217 mNm. This is the ridge of Sec. 6 realized: the free fit scatters along the amplitude-onset trade-off run by run, and its directional magnitudes sit systematically low ($|\bar{M}_{\pm}|$ smaller in 17/20 groups) — the early-onset drift of the unconstrained fit. (ii) *The deliverable survives either way*: against the load-cell truth the CoM error is statistically indistinguishable (COSH median 1.05 / RMS 1.64 / max 3.36 mm; NLS 0.86 / 1.72 / 3.28 mm), because the pivot-free average (8) cancels the largely common-mode early-onset bias between the $+$ and $-$ directions, and averaging 5–7 runs per direction suppresses the scatter. (iii) *The calibration is what converts per-run estimates from indicative to trustworthy*: with COSH a *single* run already sits within a few percent of the directional mean, which is what enables run-count reduction and per-run gating diagnostics; the NLS estimator needs the full ensemble to average its ridge scatter away.

9 Reproducibility

Every number in this document is reproducible from the repository: `critical_value_getter_piecewise.py` (pipeline, gates, caps, calibration, Algorithm 1); `analysis/ramp_quality.py` (gates); `analysis/sensitivity.py`

case	axis	λ_{truth}	COSH			NLS		
		[mm]	M_{ff} [N m]	$\hat{\lambda}$ [mm]	err	M_{ff} [N m]	$\hat{\lambda}$ [mm]	err
1	x	-2.90	-0.052	-1.73	1.17	-0.062	-2.05	0.85
1	y	-11.45	+0.329	-10.93	0.52	+0.328	-10.91	0.54
2	x	-14.29	-0.481	-15.22	0.93	-0.472	-14.95	0.66
2	y	-9.90	+0.327	-10.36	0.46	+0.306	-9.67	0.23
3	x	-5.26	-0.272	-8.62	3.36	-0.270	-8.54	3.28
3	y	+3.14	-0.023	+0.73	2.41	-0.033	+1.06	2.08
4	x	+6.67	+0.166	+5.26	1.41	+0.156	+4.94	1.73
4	y	+2.40	-0.005	+0.16	2.24	+0.020	-0.64	3.04
5	x	+10.91	+0.338	+10.70	0.21	+0.317	+10.05	0.86
5	y	-10.89	+0.338	-10.70	0.19	+0.332	-10.50	0.39
median / RMS / max err [mm]			1.05 / 1.64 / 3.36			0.86 / 1.72 / 3.28		

Table 2: Deliverables against the load-cell ground truth (manuscript Table 7): pivot-free offset (8) and the CoM offset $\hat{\lambda}$, per-case weight (3.066 kg for case 1, 3.220 kg otherwise). One row per case-axis; the two estimators sit side by side.

case	axis	COSH				NLS			
		$\bar{M}_-$ [N m]	$\bar{M}_+$ [N m]	CV $_-$ [%]	CV $_+$ [%]	$\bar{M}_-$ [N m]	$\bar{M}_+$ [N m]	CV $_-$ [%]	CV $_+$ [%]
1	x	-1.054	0.950	3.8	5.7	-1.018	0.894	6.4	7.7
1	y	-0.422	1.079	2.4	2.2	-0.399	1.056	8.8	7.2
2	x	-1.829	0.867	1.2	1.5	-1.792	0.848	2.5	4.2
2	y	-0.666	1.320	1.2	1.6	-0.637	1.248	4.9	5.8
3	x	-1.542	0.997	1.7	2.4	-1.573	1.033	2.4	4.1
3	y	-0.983	0.937	1.7	2.6	-0.990	0.923	3.2	7.0
4	x	-1.136	1.468	3.4	2.2	-1.151	1.463	3.8	3.5
4	y	-0.998	0.988	2.9	3.7	-0.928	0.969	10.3	6.1
5	x	-1.019	1.695	3.9	2.1	-0.989	1.624	7.2	2.6
5	y	-0.635	1.310	6.5	2.9	-0.577	1.241	15.2	5.4

Table 3: Directional critical-moment means and their ramp-rate dispersion. Bold marks the lower CV of each pair — COSH in all 20/20 groups.

($\pm 20\%$ robustness); `analysis/cosh_fidelity.py` (prediction-test residuals); `analysis/excitation_angle_design.py` (rate/cutoff design); `analysis/ge_linearity.py`, `ge_trajectory.py` (ground-effect bracket); `analysis/tiltcap_ablation.py` (window ablation); `analysis/baseline_stat_ablation.py` (median/mean/zero baseline); `analysis/bias_drift.py` (in-window baseline drift); `analysis/nls_comparison.py` (Table 2).