

How (95)–(97) are obtained

a step-by-step companion to Sec. VI-E

every step below is checked by `analysis/vie_geometric_box.py`

0. What is being bounded, and why the rate

The estimator does not see the tilt angle. It fits the two-segment model to the *rate*

$$\omega_{\text{nom}}(\tau) = C_1 (\cosh C_2 \tau - 1), \quad C_1 = \frac{\dot{M}}{Wz_{CoM}}, \quad C_2 = \sqrt{\frac{Wz_{CoM}}{J_P}}, \quad (\text{D.1})$$

reconstructed from the gyro, and reports the onset as the argument that minimises the residual of that fit. Two questions follow, and (97) answers both:

- (i) how far can the *true* rate depart from (D.1) before the fit stops describing the data? — the quantity $|e_\omega|_{\text{end}}/\omega_{\text{nom, end}}$;
- (ii) how far does that departure move the *reported onset moment*? — the quantity $|\Delta M_{\text{crit}}|$.

These are different questions with different answers, which is why the two halves of (97) carry different constants ($\frac{1}{2}$ and 1 in the first, $\frac{1}{7}$ and $\frac{1}{5}$ in the second). Section 5 answers (i) and Section 6 answers (ii). If only one line of this document is read, let it be that one.

The departure obeys the Duhamel integral (88),

$$e_\omega(\tau) = \int_0^\tau \dot{h}_\varphi(\tau - s) \rho(s) ds, \quad \dot{h}_\varphi(u) = \frac{\cosh C_2 u}{J_P}, \quad (\text{D.2})$$

with ρ the part of the true forcing that the estimator's span $\{1, \tau, \delta\varphi\}$ cannot absorb — the gravity remainder ρ_φ and the bilinear ground-effect coupling ρ_{GE} of (89). Everything below is about turning (D.2) into a number.

1. The crude bound, and why it is not enough

Put $|\rho(s)| \leq \rho_{\text{max}}$ inside (D.2) and pull it out:

$$|e_\omega(\tau_{\text{end}})| \leq \rho_{\text{max}} \int_0^{\tau_{\text{end}}} \dot{h}_\varphi(\tau_{\text{end}} - s) ds = \rho_{\text{max}} \int_0^{\tau_{\text{end}}} \frac{\cosh C_2 u}{J_P} du = \frac{\rho_{\text{max}} \sinh x}{J_P C_2}, \quad x \triangleq C_2 \tau_{\text{end}}. \quad (\text{D.3})$$

The substitution $u = \tau_{\text{end}} - s$ was used in the middle step; the integral of cosh is sinh. Using $J_P = Wz_{CoM}/C_2^2$, i.e. $J_P C_2 = Wz_{CoM}/C_2$, the inertia disappears: $|e_\omega(\tau_{\text{end}})| \leq \rho_{\text{max}} C_2 \sinh x / Wz_{CoM}$. This is (90).

(D.3) is honest but blunt. It treats ρ as though it sat at its maximum for the whole window, whereas both channels start at zero at the onset and grow like τ^6 and τ^4 . Evaluated over the geometric box of Sec. VI-D it returns a deviation of 446% at the slowest ramp — bigger than the signal it is supposed to bound. Sections 2–4 recover the lost factor.

2. Chebyshev's integral inequality

Statement. Let f, g be integrable on $[a, b]$. If they are *oppositely* monotone (one non-decreasing, the other non-increasing), then

$$\int_a^b fg \leq \frac{1}{b-a} \int_a^b f \int_a^b g. \quad (\text{D.4})$$

If they are monotone in the *same* direction, the inequality reverses. (Textbooks usually print the same-direction version, with $\geq$; the one needed here is the reversed one, so the direction is worth checking rather than remembering.)

Proof. Opposite monotonicity means $(f(u) - f(v))(g(u) - g(v)) \leq 0$ for every $u, v \in [a, b]$: when $u > v$ one bracket is ≥ 0 and the other ≤ 0 . Integrate that product over the square $[a, b]^2$ and expand the four terms,

$$\underbrace{\int \int f(u)g(u)}_{(b-a) \int fg} - \underbrace{\int \int f(u)g(v)}_{\int f \int g} - \underbrace{\int \int f(v)g(u)}_{\int f \int g} + \underbrace{\int \int f(v)g(v)}_{(b-a) \int fg} = 2(b-a) \int fg - 2 \int f \int g \leq 0,$$

which is (D.4). $\square$

Application to (D.2). Take $[a, b] = [0, \tau_{\text{end}}]$ with $f(s) = \rho(s)$ and $g(s) = \dot{h}_\varphi(\tau_{\text{end}} - s)$.

- f is non-decreasing: both channels vanish at the onset and grow monotonically over the window ($\rho_\varphi \propto \varphi^2$ with φ rising, $\rho_{GE} \propto \tau\varphi$ likewise). Equivalently, both are power series in u with non-negative coefficients, so $f' \geq 0$ term by term — the same sign pattern that Sec. 3 uses for a different purpose. This is the one place the argument needs the channels to increase; the reduction factors of Sec. 3 do not.
- g is non-increasing *in s* : as s grows the argument $\tau_{\text{end}} - s$ shrinks and cosh decreases. Note the kernel is increasing in its own argument; it is the reflection $s \mapsto \tau_{\text{end}} - s$ that makes it decreasing in s .

They are oppositely monotone, so (D.4) applies with $\leq$, which is the useful direction:

$$e_\omega(\tau_{\text{end}}) \leq \underbrace{\frac{1}{\tau_{\text{end}}} \int_0^{\tau_{\text{end}}} \rho}_{\bar{\rho}} \cdot \int_0^{\tau_{\text{end}}} \dot{h}_\varphi(\tau_{\text{end}} - s) ds = \frac{\bar{\rho} \sinh x}{J_P C_2}. \quad (\text{D.5})$$

(D.5) is (D.3) with ρ_{max} replaced by the *time average* $\bar{\rho}$. Nothing else changed. All that remains is to compute $\bar{\rho}/\rho_{\text{max}}$ for the two channels.

3. The reduction factors R_φ and R_{GE} : (95)

Both channels have a known shape, because φ does. From (D.1), $\varphi_{\text{nom}}(\tau) = (C_1/C_2)\Lambda(C_2\tau)$ with $\Lambda(u) = \sinh u - u$. Hence, writing $u = C_2\tau$,

$$\rho_\varphi(\tau) = \frac{1}{2}G''\varphi^2 = A\Lambda(u)^2, \quad \rho_{GE}(\tau) = \beta_M \dot{M}\tau\varphi = B u \Lambda(u),$$

with A, B constants whose values never matter — they cancel in the ratio below.

Gravity channel. The maximum is at the window end, $\rho_{\varphi, \text{max}} = A\Lambda(x)^2$. The average is, substituting $u = C_2\tau$ so that $d\tau = du/C_2$ and $\tau_{\text{end}} = x/C_2$,

$$\bar{\rho}_\varphi = \frac{1}{\tau_{\text{end}}} \int_0^{\tau_{\text{end}}} A\Lambda(C_2\tau)^2 d\tau = \frac{C_2}{x} \cdot \frac{A}{C_2} \int_0^x \Lambda(u)^2 du = \frac{A}{x} \int_0^x \Lambda(u)^2 du,$$

so the constant A cancels and

$$R_\varphi(x) \triangleq \frac{\bar{\rho}_\varphi}{\rho_{\varphi,\max}} = \frac{1}{x} \frac{\int_0^x \Lambda(u)^2 du}{\Lambda(x)^2}. \quad (\text{D.6})$$

Ground-effect channel. Identically, $\rho_{GE,\max} = B x \Lambda(x)$ and

$$R_{GE}(x) = \frac{1}{x} \frac{\int_0^x u \Lambda(u) du}{x \Lambda(x)}. \quad (\text{D.7})$$

How to read $1/7$ and $1/5$. The quickest way to see where these come from — not the proof, which is below — is the order of each channel. Keeping the leading term of $\Lambda(u) = u^3/6 + u^5/120 + \dots$ makes $\Lambda^2 \simeq u^6/36$ and $u\Lambda \simeq u^4/6$, so the two channels are sixth and fourth order near the onset, and a τ^n signal averages over $[0, T]$ to $T^n/(n+1)$ against its endpoint value T^n : hence $\frac{1}{7}$ and $\frac{1}{5}$. At the other end $\Lambda \rightarrow \sinh u \simeq e^u/2$ gives $\int_0^x \Lambda^2 \simeq e^{2x}/8$ and $\int_0^x u\Lambda \simeq x e^x/2$, so $R_\varphi \simeq 1/(2x)$ and $R_{GE} \simeq 1/x$ — the origin of the $\frac{1}{2}$ and 1 that Sec. 5 needs, once $\Psi \simeq x$ multiplies them.

Neither reading is used as a proof. (D.6b) below obtains the small- x limits from l'Hôpital's rule with no expansion of f whatever, and Sec. 5 obtains the large- x ones the same way. Only the monotonicity argument that follows uses a series, and only for its sign pattern.

Why the small- x value is the supremum. Both R 's decrease throughout, and this needs no numerics. Substituting $u = xt$ in the numerator rescales the window to $[0, 1]$:

$$R(x) = \frac{1}{x} \frac{\int_0^x f(u) du}{f(x)} = \int_0^1 \frac{f(xt)}{f(x)} dt. \quad (\text{D.6a})$$

Let

$$n_{\text{eff}}(u) \triangleq \frac{u f'(u)}{f(u)} = \frac{d \log f}{d \log u}$$

be the *local order* of f . The small- x limit of R follows from it by one application of l'Hôpital's rule, with no expansion of f : in the first form of (D.6a) numerator and denominator both vanish at $x = 0$, so

$$\lim_{x \rightarrow 0} R(x) = \lim_{x \rightarrow 0} \frac{\int_0^x f(u) du}{x f(x)} = \lim_{x \rightarrow 0} \frac{f(x)}{f(x) + x f'(x)} = \frac{1}{1 + n_{\text{eff}}(0^+)}, \quad (\text{D.6b})$$

the last step dividing through by f . Only the limit of the local order is needed, never the coefficients of f . Differentiating the integrand of (D.6a) logarithmically,

$$\frac{\partial}{\partial x} \log \frac{f(xt)}{f(x)} = t \frac{f'(xt)}{f(xt)} - \frac{f'(x)}{f(x)} = \frac{n_{\text{eff}}(xt) - n_{\text{eff}}(x)}{x}.$$

For $t \in (0, 1)$ we have $xt < x$, so if n_{eff} is increasing the bracket is ≤ 0 : every integrand of (D.6a) decreases in x , hence so does R . $\square$

Both channels satisfy the hypothesis, and with $\Lambda' = \cosh u - 1$ the local orders are expressible through the same $\Psi_\varphi(u) = u(\cosh u - 1)/(\sinh u - u) = u\Lambda'/\Lambda$ that Sec. 4 discarded as the wrong normaliser:

$$f = \Lambda^2 : \quad n_{\text{eff}} = 2\Psi_\varphi(u) \text{ from } 6 \text{ to } \infty; \quad f = u\Lambda : \quad n_{\text{eff}} = 1 + \Psi_\varphi(u) \text{ from } 4 \text{ to } \infty.$$

Why n_{eff} increases. What is required here is easy to understate, so it is worth naming precisely. Since $n_{\text{eff}} = d \log f / d \log u$, the claim

$$\frac{dn_{\text{eff}}}{d \log u} = \frac{d^2(\log f)}{d(\log u)^2} \geq 0$$

says that $\log f$ is *convex* in $\log u$ — that the slope of f on log–log axes grows. Monotonicity of f gives only $n_{\text{eff}} \geq 0$, which is the slope’s *sign* and nothing more, and the two properties are independent: $f(u) = 1 - e^{-u}$, $u/(1+u)$ and $\log(1+u)$ are all strictly increasing and all have n_{eff} strictly *decreasing*, hence R increasing. A saturating signal has a higher average relative to its endpoint the longer the window, not a lower one. The two channels here behave the opposite way not because they rise but because they *stiffen*.

Convexity does hold, and one place a power series is used — only for its sign pattern, with no coefficient evaluated — establishes it.

Lemma (log–log convexity). Let $f(u) = \sum_{n \geq 0} a_n u^n$ converge for all $u \geq 0$ with every $a_n \geq 0$, at least two of them non-zero. Then $n_{\text{eff}} = uf'/f$ is strictly increasing on $(0, \infty)$.

The second condition is what makes $f > 0$ for $u > 0$, so that dividing by f below is legitimate, and what makes the conclusion strict. No monotonicity of f is assumed — it would be redundant, since $f' = \sum_{n \geq 1} n a_n u^{n-1} \geq 0$ term by term, and it is not used in the proof.

Proof. Differentiating n_{eff} ,

$$u \frac{dn_{\text{eff}}}{du} = \frac{uf'}{f} + \frac{u^2 f''}{f} - \left(\frac{uf'}{f} \right)^2 = \frac{S_0 S_2 - S_1^2}{S_0^2}, \quad S_k \triangleq \sum_n n^k a_n u^n,$$

since $uf' = S_1$ and $u^2 f'' = S_2 - S_1$, so the first two terms combine to S_2/S_0 . Every S_k runs from $n = 0$, including the two that come from differentiating: f' starts at $n = 1$ and f'' at $n = 2$, but the coefficients n and $n(n-1) = n^2 - n$ vanish on precisely the terms those shifts drop, so extending all three sums to the same index set neither adds nor loses anything. It is worth doing, because Cauchy–Schwarz is about to treat them as sequences over one index set, and those sequences are infinite: Λ , and with it Λ^2 and $u\Lambda$, has infinitely many non-zero coefficients. That is what the hypothesis of convergence for every $u \geq 0$ is for. A power series converging at some $u_0 > 0$ converges absolutely for $|u| < u_0$, so convergence at every $u \geq 0$ makes the radius infinite, and the termwise derivatives $\sum n a_n u^{n-1}$ and $\sum n(n-1) a_n u^{n-2}$ then converge on the same disc. Hence f'' exists as written and S_1, S_2 are finite — the two sequences lie in ℓ^2 , which is what Cauchy–Schwarz for infinite sums requires.

Cauchy–Schwarz applied to $\{\sqrt{a_n u^n}\}$ and $\{n\sqrt{a_n u^n}\}$ — the square roots real because $a_n u^n \geq 0$, which is the only place the sign pattern is needed — gives $S_1^2 \leq S_0 S_2$. Hence $dn_{\text{eff}}/du \geq 0$, with equality only when the two sequences are proportional, i.e. only when a single a_n is non-zero. $\square$

(Dividing through by S_0 turns S_k/S_0 into the moments of the weights $a_n u^n/f$ and the numerator into their variance, which is a shorter way to say the same thing; nothing below depends on reading it that way.)

Both kernels qualify: $\Lambda = \sum_{k \geq 1} u^{2k+1}/(2k+1)!$ has non-negative coefficients, and so do the product Λ^2 and the shift $u\Lambda$; each has infinitely many non-zero, so certainly two. Hence n_{eff} is strictly increasing for both channels and R is strictly decreasing. Reading the $a_n u^n/f$ as weights also returns the origin values, since as $u \rightarrow 0$ they concentrate on the lowest exponent present, 6 and 4 — the same numbers the l’Hôpital route below obtains without any series.

The same fact as a closed-form derivative. An appendix that quotes dR/dx explicitly needs those expressions signed, and the argument above does it with no further machinery — in particular without the twelve-stage differentiation that signing them directly would cost. Write $A(x) = \int_0^x f$, so $R = A/(xf)$. Then, using $A' = f$ and $xf' = n_{\text{eff}} f$,

$$R'(x) = \frac{A' x f - A(f + x f')}{x^2 f^2} = \frac{x f(x) - (1 + n_{\text{eff}}(x)) A(x)}{x^2 f(x)}. \quad (\text{D.7a})$$

The sign turns on whether A exceeds $xf/(1 + n_{\text{eff}})$, and it does, for exactly the reason n_{eff}

increases. Since $d \log f / d \log u = n_{\text{eff}}$ and $n_{\text{eff}}(t) \leq n_{\text{eff}}(x)$ for $t \leq x$,

$$\log \frac{f(x)}{f(u)} = \int_u^x n_{\text{eff}}(t) \frac{dt}{t} \leq n_{\text{eff}}(x) \log \frac{x}{u} \implies f(u) \geq f(x) \left(\frac{u}{x} \right)^{n_{\text{eff}}(x)},$$

and integrating that minorant over $[0, x]$,

$$A(x) \geq f(x) \int_0^x \left(\frac{u}{x} \right)^{n_{\text{eff}}(x)} du = \frac{x f(x)}{1 + n_{\text{eff}}(x)},$$

strictly, since n_{eff} is not constant. The numerator of (D.7a) is therefore negative and $R' < 0$. $\square$

In words: f is at least as steep as the power law of its local order at x , so its average over the window is at least the $1/(n+1)$ that power law would give — and $1/(n+1)$ is precisely the value R would need for R' to vanish.

Substituting $n_{\text{eff}} = 2x\Lambda'/\Lambda$ for $f = \Lambda^2$, and $n_{\text{eff}} = (\Lambda + x\Lambda')/\Lambda$ for $f = u\Lambda$, and clearing denominators puts (D.7a) in the form an appendix would quote:

$$\frac{dR_\varphi}{dx} = \frac{x\Lambda^3 - (\Lambda + 2x\Lambda') \int_0^x \Lambda^2 du}{x^2 \Lambda^3}, \quad \frac{dR_{GE}}{dx} = \frac{x^2 \Lambda^2 - (2\Lambda + x\Lambda') \int_0^x u\Lambda du}{x^3 \Lambda^2}, \quad (\text{D.7b})$$

both negative for $x > 0$. The factors of x inside the brackets are load-bearing: with them the numerators cancel to order x^{12} and x^{10} at the origin, which is what leaves R' finite there; drop either and the expression diverges as x^{-2} instead. Both forms are checked against a five-point difference of R by `analysis/closed_form_check.py`.

The origin values without the series. The route above needs the coefficients' sign pattern; this one needs nothing, and is worth recording because it settles the same two numbers independently. The value of Ψ_φ at the origin is l'Hôpital's rule, applied three times. Writing $\Psi_\varphi = N/D$ with $N = u \cosh u - u$ and $D = \sinh u - u$, successive differentiation gives the four pairs

$$\frac{u \cosh u - u}{\sinh u - u}, \quad \frac{\cosh u + u \sinh u - 1}{\cosh u - 1}, \quad \frac{2 \sinh u + u \cosh u}{\sinh u}, \quad \frac{3 \cosh u + u \sinh u}{\cosh u},$$

of which the first three are $0/0$ at $u = 0$ while the fourth is $3/1$. Hence $\Psi_\varphi(0^+) = 3$, so $n_{\text{eff}}(0^+)$ is 6 for Λ^2 and 4 for $u\Lambda$, and (D.6b) returns $\frac{1}{7}$ and $\frac{1}{5}$. The mechanism is physical: Λ behaves as u^3 near the onset and grows exponentially later, so the larger the window the more sharply both channels are concentrated at its end, and the smaller their average is relative to the endpoint value. The small- x values are therefore suprema:

$$\bar{\rho} = R_\varphi(x) \rho_{\varphi, \text{max}} + R_{GE}(x) \rho_{GE, \text{max}}, \quad R_\varphi \leq \frac{1}{7}, \quad R_{GE} \leq \frac{1}{5}. \quad (95)$$

The suprema are approached but never attained: for $x > 0$ the inequalities are strict. Over the flown $x \in [1.79, 5.20]$ the factors run 0.089–0.133 and 0.145–0.191, so quoting $\frac{1}{7}$ and $\frac{1}{5}$ gives away 10–60% — while gaining a factor of 5 to 11 over (D.3).

4. Normalising on the rate: (96)

Divide (D.5) by the nominal rate at the same instant, $\omega_{\text{nom, end}} = C_1(\cosh x - 1)$ with $C_1 = \dot{M}/Wz_{CoM}$, and use $J_P C_2 = Wz_{CoM}/C_2$:

$$\frac{|e_\omega|_{\text{end}}}{\omega_{\text{nom, end}}} \leq \frac{\bar{\rho} C_2 \sinh x / Wz_{CoM}}{(\dot{M}/Wz_{CoM})(\cosh x - 1)} = \frac{\bar{\rho} C_2}{\dot{M}} \cdot \frac{\sinh x}{\cosh x - 1}.$$

Wz_{CoM} has cancelled. Now eliminate $\dot{M}$ in favour of the moment increment the window actually spans. Since $\tau_{\text{end}} = x/C_2$,

$$\Delta M_{\text{win}} = \dot{M} \tau_{\text{end}} = \frac{\dot{M} x}{C_2} \iff \frac{C_2}{\dot{M}} = \frac{x}{\Delta M_{\text{win}}},$$

and substituting,

$$\frac{|e_\omega|_{\text{end}}}{\omega_{\text{nom, end}}} \leq \frac{\Psi(x) \bar{\rho}}{\Delta M_{\text{win}}}, \quad \Psi(x) = \frac{x \sinh x}{\cosh x - 1}. \quad (96)$$

The half-angle identities $\cosh x - 1 = 2 \sinh^2(x/2)$ and $\sinh x = 2 \sinh(x/2) \cosh(x/2)$ collapse the quotient:

$$\Psi(x) = x \frac{2 \sinh(x/2) \cosh(x/2)}{2 \sinh^2(x/2)} = x \coth \frac{x}{2} \geq 2.$$

Ψ rises from 2 at $x \rightarrow 0$ to x as $x \rightarrow \infty$. (Had the excursion been used as the normaliser instead of the rate, the factor would have been $x(\cosh x - 1)/(\sinh x - x)$, rising from 3 — a different function, and the wrong one, since the estimator never looks at φ .)

5. Why $\Psi \bar{\rho}$ stays bounded: first half of (97)

Ψ grows with x and the R 's shrink with it, and the two products converge to $\frac{1}{2}$ and 1. Both bounds are exact, and neither needs the series: one corollary of the mean value theorem settles both, three applications for the first and one for the second.

The claims in integral form. Since $\Psi = x \coth(x/2)$ and $\rho_{\varphi, \text{max}} = A\Lambda(x)^2$,

$$\Psi R_\varphi = x \coth \frac{x}{2} \cdot \frac{\int_0^x \Lambda^2 du}{x \Lambda(x)^2} = \coth \frac{x}{2} \frac{\int_0^x \Lambda^2 du}{\Lambda(x)^2},$$

so $\Psi R_\varphi \leq \frac{1}{2}$ is exactly

$$\int_0^x \Lambda(u)^2 du \leq \frac{1}{2} \tanh \frac{x}{2} \Lambda(x)^2, \quad (\text{D.8})$$

and the same reduction turns $\Psi R_{GE} \leq 1$ into

$$\int_0^x u \Lambda(u) du \leq x \tanh \frac{x}{2} \Lambda(x). \quad (\text{D.9})$$

The half-angle tangent is not an accident: $\tanh(x/2) = (\cosh x - 1)/\sinh x = \Lambda'(x)/\sinh x$, so both statements compare an integral of Λ against Λ and its own derivative.

Where $\frac{1}{2}$ and 1 come from. Both are the $x \rightarrow \infty$ limits, and each is one application of l'Hôpital's rule. Written as above, both products have the form $\coth(x/2) \int_0^x f/f(x)$ — with $f = \Lambda^2$ for the gravity channel and $f = u\Lambda$ for the ground-effect channel — and $\coth(x/2) \rightarrow 1$. In the remaining ratio numerator and denominator both diverge, so

$$\lim_{x \rightarrow \infty} \frac{\int_0^x f(u) du}{f(x)} = \lim_{x \rightarrow \infty} \frac{f(x)}{f'(x)} = \left(\lim_{x \rightarrow \infty} \frac{f'}{f} \right)^{-1},$$

and the two logarithmic derivatives are elementary:

$$\left. \frac{f'}{f} \right|_{f=\Lambda^2} = \frac{2\Lambda'}{\Lambda} = \frac{2(\cosh x - 1)}{\sinh x - x} \rightarrow 2, \quad \left. \frac{f'}{f} \right|_{f=u\Lambda} = \frac{1}{u} + \frac{\Lambda'}{\Lambda} \rightarrow 1,$$

giving $\frac{1}{2}$ and 1. Each constant is thus the reciprocal of its kernel's exponential growth rate, which is what $\int_0^x e^{ku} du = e^{kx}/k$ says: the gravity channel gets the smaller constant because its kernel, $\Lambda^2 \sim e^{2x}/4$, is the square of the other's, $u\Lambda \sim u e^x/2$.

The only theorem used. Everything below rests on one corollary of the mean value theorem, applied four times in all — three for (D.8) and once for (D.9) — and nowhere supplemented.

Lemma (a corollary of the mean value theorem). If g is continuous on $[0, b]$, differentiable on $(0, b)$ with $g' \geq 0$ there, and $g(0) = 0$, then $g \geq 0$ on $[0, b]$. *Proof:* for $0 \leq s < t \leq b$ the mean

value theorem gives $g(t) - g(s) = g'(\xi)(t - s) \geq 0$, so g is non-decreasing; taking $s = 0$ gives $g(t) \geq g(0) = 0$. $\square$ (The symbol is g , not the f of Sec. 3: this lemma is applied to ϕ' , ϕ and D , never to a channel kernel.)

Note what is *not* used. No limit as $x \rightarrow \infty$ enters: the constant $\frac{1}{2}$ comes from the algebra of (D.8), and the inequality is established at each x separately. No monotonicity of ΨR_φ is assumed or proved. The limits quoted above serve only to show the constants cannot be improved.

Proof of (D.8). Let $D(x)$ be the right side minus the left, so that (D.8) is $D \geq 0$. Every quantity divided by below $-\coth(x/2)$, Λ^2 , Λ , $\cosh^2(x/2)$ — is strictly positive for $x > 0$, so each step is an equivalence.

$D(0) = 0$, since $\tanh 0 = \Lambda(0) = \int_0^0 = 0$. Differentiating and factoring out $\Lambda > 0$,

$$D'(x) = \Lambda(x) B(x), \quad B(x) = \frac{1}{4} \operatorname{sech}^2 \frac{x}{2} \Lambda + \tanh \frac{x}{2} \Lambda' - \Lambda,$$

so $D' \geq 0$ iff $B \geq 0$. Write $y = x/2$, $S = \sinh y$, $C = \cosh y$, giving $\Lambda = 2SC - 2y$ and $\Lambda' = 2S^2$. Multiplying B by $2C^2 > 0$,

$$2C^2 B = (SC - y) + 4S^3 C - 4SC^3 + 4yC^2,$$

and since $S^3 C - SC^3 = SC(S^2 - C^2) = -SC$ the two cubic terms collapse to $-4SC$, leaving $y(4C^2 - 1) - 3SC = y(3 + 4S^2) - 3SC$. Now $4S^2 = 2 \cosh x - 2$ and $2SC = \sinh x$ and $y = x/2$, so

$$B(x) = \frac{\phi(x)}{4 \cosh^2(x/2)}, \quad \phi(x) \triangleq x + 2x \cosh x - 3 \sinh x,$$

and $B \geq 0$ iff $\phi \geq 0$. Finally

$$\phi'(x) = 1 - \cosh x + 2x \sinh x, \quad \phi''(x) = \sinh x + 2x \cosh x.$$

First application. $g = \phi'$: both terms of $\phi'' = \sinh x + 2x \cosh x$ are non-negative for $x \geq 0$, and $\phi'(0) = 1 - 1 + 0 = 0$. The Lemma gives $\phi' \geq 0$.

Second application. $g = \phi$: $\phi' \geq 0$ by the first application and $\phi(0) = 0$. The Lemma gives $\phi \geq 0$, hence $B \geq 0$, hence $D' = \Lambda B \geq 0$.

Third application. $g = D$: $D' \geq 0$ by the second application and $D(0) = 0$. The Lemma gives $D \geq 0$, which is (D.8). $\square$

Proof of (D.9). The same Lemma settles the ground-effect channel, and here it is needed only once. Let E be the right side of (D.9) minus the left,

$$E(x) = x \tanh \frac{x}{2} \Lambda(x) - \int_0^x u \Lambda(u) du, \quad E(0) = 0.$$

Differentiating — the integral contributes $x\Lambda(x)$ — and using $\operatorname{sech}^2(x/2) = 2/(\cosh x + 1)$ and $\tanh(x/2) = \sinh x/(\cosh x + 1)$ to put everything over $\cosh x + 1$, the hyperbolic terms collapse to

$$E'(x) = \frac{\sinh^2 x - 2x \sinh x + x^2 \cosh x}{\cosh x + 1} = \frac{\Lambda(x)^2 + x^2 \Lambda'(x)}{\cosh x + 1}, \quad (\text{D.10})$$

the second equality being the identity $(\sinh x - x)^2 + x^2(\cosh x - 1) = \sinh^2 x - 2x \sinh x + x^2 \cosh x$. The numerator is a sum of two non-negative terms and the denominator is positive, so $E' \geq 0$; the Lemma gives $E \geq 0$, which is (D.9). $\square$

No series is needed for either, and (D.9) costs one application of the Lemma against three for (D.8) — the completed square in (D.10) does the work that ϕ did there.

No monotonicity of the products is claimed or needed — only that they do not exceed their limits, which is what (D.8) and (D.9) say. Expanding $\bar{\rho}$ by (95),

$$\Psi \bar{\rho} = (\Psi R_\varphi) \rho_{\varphi, \max} + (\Psi R_{GE}) \rho_{GE, \max} \leq \frac{1}{2} \rho_{\varphi, \max} + \rho_{GE, \max},$$

which is the left half of (97). *x has disappeared* — and with it C_2 , J_P and Wz_{CoM} . This is the whole point of the exercise: the bound no longer depends on the window length, so it does not have to be solved for, only bounded well enough to keep $\rho_{GE,\max}$ finite (which is what (92)–(93) do).

6. From the rate deviation to the threshold: second half of (97)

The reported quantity is not e_ω but the onset moment, so the deviation must be pushed one step further.

How the onset responds to a perturbation. The fit is two-segment and *both* segments are in the cost — a constant baseline before the onset, the closed-form rise after it, sharing that baseline:

$$J(t_c) = \int_{t_a}^{t_c} (y - C_0)^2 dt + \int_{t_c}^{t_{\text{end}}} \left(y - C_0 - C_1 (\cosh C_2(t - t_c) - 1) \right)^2 dt,$$

with C_0 the pre-onset baseline, refitted at each candidate t_c from the samples to its left. In the shifted variable $\tau = t - t_c$ the post-onset window is $[0, T]$ with $T = t_{\text{end}} - t_c = \tau_{\text{end}}$, the window Sec. 2 averages over, so the $\bar{\rho}$ below is the same $\bar{\rho}$.

Write the data as the true pre-onset level, the nominal response from the true onset, and a deviation,

$$y(t) = b + \hat{\omega}(t; t_c^*) + e_\omega(t), \quad t_c = t_c^* + \delta t_c,$$

the perturbation being ρ itself and not its average. Carrying b rather than setting it to zero is what later gives C_0 something to estimate; without it C_0 would enter the cost and leave the stationarity condition unaccounted for.

One point about that definition is worth making explicitly, because getting it wrong invalidates everything after it. *e_ω is anchored at the true onset t_c^* , not at the candidate.* It is the difference between the recorded rate and the nominal solution of the true problem, so once the run is recorded it is a fixed function of t and

$$\frac{\partial e_\omega}{\partial t_c} = 0 \quad \text{identically.}$$

When the residual is written $e_\omega - \delta t_c \chi - C$ below, every appearance of the optimisation variable is in δt_c and in C ; none is in e_ω . Reading e_ω instead as “the residual at the current candidate” would make it depend on t_c and would count the same dependence twice.

Name the sensitivity

$$\chi(t) \triangleq \left. \frac{\partial \hat{\omega}(t; t_c)}{\partial t_c} \right|_{t_c=t_c^*} = -C_1 C_2 \sinh(C_2(t - t_c^*)).$$

It needs a symbol of its own, and not φ : that is the tilt angle everywhere else, and $\langle e_\omega, \varphi \rangle$ would read as a rate correlated against an angle. To first order $\hat{\omega}(t; t_c^* + \delta t_c) \simeq \hat{\omega}(t; t_c^*) + \delta t_c \chi$, so the post-onset residual is $e_\omega - \delta t_c \chi - (C_0 - b)$: the baseline enters only through its *error*, a point taken up below. Setting that error aside for the moment,

$$\frac{\partial J}{\partial t_c} = \underbrace{(y(t_c) - C_0)^2 - (y(t_c) - C_0)^2}_{= 0, \text{ the two limits}} - 2 \int_{t_c}^{t_{\text{end}}} (e_\omega - \delta t_c \chi) \chi dt - 2 \frac{dC_0}{dt_c} \int r = 0.$$

t_c is the upper limit of one integral and the lower limit of the other, so it produces two boundary terms, and they cancel *identically*: at $\tau = 0$ the rise contributes $C_1(\cosh 0 - 1) = 0$, so the two-segment model is continuous at the onset by construction and both terms are $(y(t_c) - C_0)^2$ with opposite signs. Nothing is discarded and no order argument is needed. This is worth stating, because the first thing to check about a fit whose window endpoint is also its optimisation variable is whether the Leibniz terms were quietly dropped.

The baseline still couples the two segments, since C_0 depends on t_c through the samples it is fitted to. Its contribution to the pre-onset integral is $(dC_0/dt_c) \int_{t_a}^{t_c} (y - C_0)$, which vanishes — for the least-squares constant by its normal equation, and for the median the code actually uses because a median is piecewise constant in the sample set, so $dC_0/dt_c = 0$ almost everywhere. What could survive is the same drift entering the post-onset residual, but $dC_0/dt_c = (y(t_c) - C_0)/(t_c - t_a)$ for the mean, and at the onset the rate is still the baseline, so it vanishes there too. Measured across the 140 runs, C_0 moves by under 7.4×10^{-4} rad/s when the candidate onset shifts by one sample. Dropping that term and solving,

$$\delta t_c = \frac{\langle e_\omega, \chi \rangle}{\|\chi\|^2}, \quad \langle a, b \rangle \triangleq \int_{t_c^*}^{t_{\text{end}}} a(t) b(t) dt, \quad (\text{D.11})$$

the division legitimate because $\|\chi\|^2 > 0$ — $\sinh(C_2\tau)$ is not identically zero on a window of positive length.

Where the second derivative of $\hat{\omega}$ went. Only χ appears in the stationarity condition because $\partial J/\partial t_c$ differentiates the residual once; $\partial_{t_c}\chi$ belongs to the *second* derivative of J , and (D.11) is the Gauss–Newton form that drops it. Writing the one-parameter Newton step exactly,

$$J'(t_c^*) = -2\langle e_\omega, \chi \rangle, \quad J''(t_c^*) = 2\|\chi\|^2 - 2\langle e_\omega, \partial_{t_c}\chi \rangle,$$

$$\delta t_c = -\frac{J'}{J''} = \frac{\langle e_\omega, \chi \rangle}{\|\chi\|^2 - \langle e_\omega, \partial_{t_c}\chi \rangle}, \quad \partial_{t_c}\chi = C_1 C_2^2 \cosh(C_2\tau).$$

The neglected term is the residual times a curvature, the one nonlinear least squares always discards. It is legitimate here for a reason worth stating plainly: *the expansion parameter is e_ω , not δt_c* , and δt_c is itself first order in e_ω by (D.11). So $\langle e_\omega, \chi \rangle$ and $\delta t_c \|\chi\|^2$ are both $\mathcal{O}(e_\omega)$ while $\delta t_c \langle e_\omega, \partial_{t_c}\chi \rangle$ is $\mathcal{O}(e_\omega^2)$ — every term carrying a derivative of χ picks up an extra factor of the deviation, including the $(\delta t_c)^2 \partial_{t_c}\chi$ left over from expanding the residual itself.

Solving for δt_c directly. Once the residual is linearised, J is an *exact* quadratic in δt_c , so it is cleaner to differentiate in that variable and solve for the value that zeroes the derivative. Since $t_c = t_c^* + \delta t_c$ the operator is the same, but the two approximations then sit in separate places: linearising the residual is the only one, and the normal equation that follows is exact for the quadratic it is applied to.

What the baseline has to satisfy, and what it does not. Keeping C_0 inside that square matters. Differentiating $(e_\omega - \delta t_c \chi - C_0)^2$ and solving gives

$$\delta t_c = \frac{\langle e_\omega - C_0, \chi \rangle}{\|\chi\|^2} = \frac{\langle e_\omega, \chi \rangle}{\|\chi\|^2} - C_0 \frac{\langle 1, \chi \rangle}{\|\chi\|^2}, \quad (\text{D.11a})$$

and the second term does not vanish: χ is of one sign, so $\langle 1, \chi \rangle = -C_1(\cosh x - 1) \neq 0$.

It does not follow that C_0 must be driven to zero, which is the natural but wrong reading of the onset condition. Only the *error* of C_0 against the true pre-onset level b survives: a run whose gyro sits at $b = 0.2$ rad/s is displaced by nothing at all provided $C_0 = b$, while a mismatch of 0.001 rad/s on a run with $b = 0$ displaces the onset measurably. What is required is an unbiased estimate of the pre-onset level — which is why the estimator fits C_0 rather than fixing it, and fits a median: the pre-onset gyro noise is heavy-tailed, and a candidate that overshoots leaks the rise into the segment being averaged.

Carrying (D.11a) through $\Delta M_{\text{crit}} = \dot{M} \delta t_c$ and $\|\chi\|^2 = C_1^2 C_2 [\frac{1}{4} \sinh 2x - \frac{1}{2}x]$, the amplitude cancels as everywhere else and

$$\Delta M_{\text{crit},C} = Wz_{CoM} (C_0 - b) \frac{\cosh x - 1}{C_2 [\frac{1}{4} \sinh 2x - \frac{1}{2}x]}. \quad (\text{D.11b})$$

Minimising the cost directly reproduces (D.11b) to three figures. Over the box the gain runs from 0.042 N m per rad/s at the slowest ramp to 0.351 at the fastest, the fast end being worse because the window is shorter. The baselines the sweep actually forms — medians over the excitation window up to the candidate, not over the whole quiet record before it — are 3.5 mrad/s in the median and 47.8 at worst across the 140 runs, so bounding $|C_0 - b|$ by $|C_0|$ gives 0.054 mm of offset in the median and 0.558 mm at worst.

That bound is loose, and the direct measurement says so. Rerunning the whole pipeline with C_0 pinned to zero instead of fitted, which is the largest perturbation this term admits, moves the direction-paired M_{ff} by 0.80 mN m in the median and 4.91 at worst — 0.027 mm and 0.163 mm. The pairing cancels most of what the bound forwards. Neither statistic is better: the ramp-invariance CV is 0.0287 against 0.0292, a paired t of 0.87 on 19 degrees of freedom, and 66 of the 140 runs return an identical onset index either way. The median is kept because a pinned zero has no defence against genuine pre-onset motion, not because it measures better.

Numerically, $\langle e_\omega, \partial_{t_c} \chi \rangle / \|\chi\|^2$ is 9.05×10^{-5} , 9.05×10^{-4} , 9.05×10^{-3} and 9.05×10^{-2} as the deviation grows through four decades: exactly proportional to it, as an $\mathcal{O}(e_\omega)$ relative error must be. That ratio is the whole of the discrepancy reported below. With $\Delta M_{\text{crit}} = \dot{M} \delta t_c$ and $C_1 = \dot{M} / Wz_{CoM}$ the amplitude C_1 cancels between numerator and denominator, leaving

$$\Delta M_{\text{crit}} = - \frac{Wz_{CoM} \langle e_\omega, \sinh(C_2 \tau) \rangle}{C_2 \|\sinh(C_2 \tau)\|^2}. \quad (\text{D.12})$$

This is exactly the reduction `error_budget.py` implements, and it is the manuscript's (98)–(106) written with χ named: (D.11) is its (104), the baseline term below is its (105)–(106). Minimising J numerically against a deviation of the true shape, over four decades of deviation size, gives ratios to the formula of 0.99995, 0.99953, 0.99530 and 0.95559: it is the first-order term, as claimed. Minimising the full two-segment cost and the post-onset term alone gives the same answer to five figures at every size, which is the numerical form of the two paragraphs above.

From (D.12) to (97b), step by step. The reduction is short but has three distinct moves in it, so they are separated here. Throughout, τ runs from 0 at the onset to τ_{end} at the window end, and $\langle a, b \rangle = \int_0^{\tau_{\text{end}}} ab \, d\tau$.

Step 1: write the inner product out. (D.12) reads

$$\Delta M_{\text{crit},e} = - \frac{Wz_{CoM}}{C_2} \frac{\int_0^{\tau_{\text{end}}} e_\omega(\tau) \sinh(C_2 \tau) \, d\tau}{\|\sinh(C_2 \tau)\|^2}.$$

Step 2: substitute the deviation. e_ω is not free: (D.2) gives it as the Duhamel integral of the forcing, with $\dot{h}_\varphi(u) = \cosh(C_2 u) / J_P$ from (87). Putting that in makes the numerator a double integral over the triangle $0 \leq s \leq \tau \leq \tau_{\text{end}}$,

$$\int_0^{\tau_{\text{end}}} \sinh(C_2 \tau) \left[\int_0^\tau \frac{\cosh(C_2(\tau - s))}{J_P} \rho(s) \, ds \right] d\tau.$$

Step 3: exchange the order. Fubini applies — both integrands are continuous on a bounded domain — and the same triangle read the other way has s outermost with τ running from s up to τ_{end} :

$$= \int_0^{\tau_{\text{end}}} \rho(s) \underbrace{\left[\frac{1}{J_P} \int_s^{\tau_{\text{end}}} \sinh(C_2 \tau) \cosh(C_2(\tau - s)) \, d\tau \right]}_{\text{a function of } s \text{ alone}} ds.$$

The forcing is now outside, multiplied by something that no longer depends on it. That is the whole point of the step: what began as a correlation of e_ω against a kernel has become a weighted average of ρ itself,

$$\Delta M_{\text{crit},e} = - \int_0^{\tau_{\text{end}}} \rho(s) w(s) ds, \quad w(s) = \frac{Wz_{CoM} \int_s^{\tau_{\text{end}}} \sinh(C_2\tau) \cosh(C_2(\tau-s)) d\tau}{J_P C_2 \|\sinh(C_2\tau)\|^2}. \quad (\text{D.13})$$

Step 4: three properties of the weight. All three are checked in code (`error_budget.py --check`).

- (i) $w \geq 0$: the integrand is a product of two hyperbolic functions of non-negative arguments.
- (ii) $\int_0^{\tau_{\text{end}}} w = 1$, *exactly*. Integrate the numerator of (D.13) over s and exchange the order once more — the same triangle again —

$$\int_0^{\tau_{\text{end}}} \int_s^{\tau_{\text{end}}} \sinh(C_2\tau) \cosh(C_2(\tau-s)) d\tau ds = \int_0^{\tau_{\text{end}}} \sinh(C_2\tau) \left[\int_0^\tau \cosh(C_2(\tau-s)) ds \right] d\tau.$$

The inner integral is elementary, $u = \tau - s$ turning it into $\int_0^\tau \cosh(C_2u) du = \sinh(C_2\tau)/C_2$, so the whole numerator becomes $\|\sinh(C_2\tau)\|^2/C_2$ — the very norm sitting in the denominator of (D.13). Hence

$$\int_0^{\tau_{\text{end}}} w = \frac{Wz_{CoM} \|\sinh\|^2/C_2}{J_P C_2 \|\sinh\|^2} = \frac{Wz_{CoM}}{J_P C_2^2} = 1$$

by $J_P C_2^2 = Wz_{CoM}$, the identity that removes the inertia everywhere else in this document. Neither the window length nor any constant survives. `error_budget.py --check` evaluates (D.13) at $\rho \equiv 1$ and returns -1.0005 to -1.0001 across five operating points, which is quadrature error on an identity rather than evidence for it.

- (iii) w is non-increasing in s : as s grows both the integration range and the cosh factor shrink.

Step 5: what (i) and (ii) alone give. A non-negative weight of unit mass can do no more than pick out the largest value of what it averages, so

$$|\Delta M_{\text{crit},e}| \leq \rho_{\text{max}} \int_0^{\tau_{\text{end}}} w = \rho_{\text{max}}.$$

This is already a result — a uniform unmodelled moment displaces the identified threshold one-for-one, as a consequence rather than an assumption — but it is the weak form, and over the box it is 2.70 mm of offset, larger than the accuracy actually achieved.

Step 6: what (iii) adds. The forcing rises across the window while the weight falls, so the two are oppositely monotone and Chebyshev's inequality (D.4) applies a second time, in the same direction as the first:

$$\left| \int_0^{\tau_{\text{end}}} \rho w \right| \leq \frac{1}{\tau_{\text{end}}} \int_0^{\tau_{\text{end}}} \rho \cdot \int_0^{\tau_{\text{end}}} w = \bar{\rho} \cdot 1 = \bar{\rho}.$$

The maximum has been replaced by the window average, and by (ii) no kernel gain is left over to multiply it.

Step 7: bound the average. $\bar{\rho}$ was computed at (95) as $R_\varphi \rho_{\varphi,\text{max}} + R_{GE} \rho_{GE,\text{max}}$, and the two reduction factors are bounded there by $\frac{1}{7}$ and $\frac{1}{5}$. Chaining Steps 6 and 7,

$$|\Delta M_{\text{crit},e}| \leq \bar{\rho} \leq \frac{\rho_{\varphi,\text{max}}}{7} + \frac{\rho_{GE,\text{max}}}{5}, \quad (97b)$$

the right half of (97). Note the direction of the second inequality: $\bar{\rho}$ is the average, and the right-hand side bounds it — it is not a definition of $\bar{\rho}$.

The same bound in two lines, without Fubini. Steps 1–6 build the weight in order to expose its three properties, and (iii) in particular is what licenses Chebyshev. But if only the *value* of the bound is wanted, there is a shorter route that uses nothing beyond (90) and the projection (D.11), and it is worth recording because it answers the obvious question — whether the pointwise bound on e_ω can be used directly — with a yes.

Evaluate (90) at each τ rather than only at the window end. Since ρ rises, its running average does too, so $\bar{\rho}(\tau) \leq \bar{\rho}$ and

$$|e_\omega(\tau)| \leq E(\tau) \triangleq \frac{\bar{\rho} \sinh C_2 \tau}{J_P C_2} \quad \text{for every } \tau \in [0, \tau_{\text{end}}]. \quad (\text{D.13a})$$

The onset shift is the projection (D.11), $\delta t_c = \langle e_\omega, \chi \rangle / \|\chi\|^2$ with $\chi = -C_1 C_2 \sinh C_2 \tau$, and χ is of one sign, so the absolute value may be taken inside:

$$|\delta t_c| \leq \frac{\langle E, |\chi| \rangle}{\|\chi\|^2}. \quad (\text{D.13b})$$

Now the point. E and χ are both proportional to $\sinh C_2 \tau$, so the integral $\int_0^{\tau_{\text{end}}} \sinh^2 C_2 \tau d\tau$ appears in numerator and denominator alike and cancels without ever being evaluated:

$$\frac{\langle E, |\chi| \rangle}{\|\chi\|^2} = \frac{(\bar{\rho} C_1 / J_P) \int_0^{\tau_{\text{end}}} \sinh^2 C_2 \tau d\tau}{C_1^2 C_2^2 \int_0^{\tau_{\text{end}}} \sinh^2 C_2 \tau d\tau} = \frac{\bar{\rho}}{J_P C_1 C_2^2} = \frac{\bar{\rho}}{C_1 W z_{CoM}},$$

using $J_P C_2^2 = W z_{CoM}$ at the last step. Carrying it through $\Delta M_{\text{crit},e} = \dot{M} \delta t_c$ and $C_1 = \dot{M} / W z_{CoM}$, so that $C_1 W z_{CoM} = \dot{M}$,

$$|\Delta M_{\text{crit},e}| \leq \dot{M} \cdot \frac{\bar{\rho}}{C_1 W z_{CoM}} = \bar{\rho}, \quad (\text{D.13c})$$

which is Step 6 again. Checked numerically over a range of $(C_2, J_P, C_1, \tau_{\text{end}})$ the ratio of the two is 1 to six decimals.

The two routes are the same inequality seen twice: what appears as $\int w = 1$ in the long route appears as the cancellation of $\int \sinh^2$ in the short one. The long route is still the one to keep, because it exhibits w and its monotonicity, and without those Chebyshev has nothing to act on — Step 6 would have no justification and the bound would stop at ρ_{max} , the 2.70 mm form. The short route is the check, and the answer to a reader who asks why the deviation bound cannot simply be inserted into the budget: inserted through the projection, it can, and it gives exactly $\bar{\rho}$.

The band, and what the minimiser can reach. (D.13a) is a statement about the *data*: the recorded rate, less its baseline, lies within $\pm E(\tau)$ of the nominal cosh. The natural question is whether the same is true of the *answer* — the minimiser returns a curve, and one wants to know whether that curve stays in the band the data lies in.

It does, but the reader should be told at once in what sense, because the derivation below is a *corollary of* (D.13c) and not an independent check on it. It takes $|\delta| \leq \bar{\rho} / \dot{M}$ from (D.13c) and multiplies by $C_1 C_2 \sinh$; E was itself built from the same $\bar{\rho}$, and $W z_{CoM} = J_P C_2^2$ makes the two agree exactly. The calculation could not have come out any other way, and nothing would have been learned had it come out differently. It is worth writing because it converts a statement about a number into a statement about a curve — but the load-bearing question, whether a fitted curve must lie near the truth *at all*, is not answered by it. That question is general, has a short answer, and is taken up after this.

With the amplitude pinned at $C_1 = \dot{M}/Wz_{CoM}$, the fitted curve is the nominal one with its onset moved and nothing else:

$$\hat{\omega}(\tau) = C_1(\cosh C_2(\tau - \delta) - 1), \quad \delta \triangleq \delta t_c. \quad (\text{D.13d})$$

Subtracting the nominal and applying $\cosh A - \cosh B = 2 \sinh \frac{A+B}{2} \sinh \frac{A-B}{2}$ — an identity, no expansion yet —

$$\hat{\omega}(\tau) - \omega_{\text{nom}}(\tau) = -2C_1 \sinh \frac{C_2\delta}{2} \sinh C_2\left(\tau - \frac{\delta}{2}\right). \quad (\text{D.13e})$$

Expanding the second factor and writing $g \triangleq C_2|\delta|$, so that $2 \sinh \frac{g}{2} \cosh \frac{g}{2} = \sinh g$ and $2 \sinh^2 \frac{g}{2} = \cosh g - 1$,

$$|\hat{\omega} - \omega_{\text{nom}}| \leq C_1 \sinh g \sinh C_2\tau + C_1(\cosh g - 1) \cosh C_2\tau. \quad (\text{D.13f})$$

This is exact. To first order in g the second term drops and $\sinh g \rightarrow g$, and then the chain closes on itself:

$$\begin{aligned} |\hat{\omega}(\tau) - \omega_{\text{nom}}(\tau)| &\leq C_1 C_2 |\delta| \sinh C_2\tau \leq \frac{\bar{\rho}}{C_1 Wz_{CoM}} C_1 C_2 \sinh C_2\tau \\ &= \frac{\bar{\rho} C_2}{Wz_{CoM}} \sinh C_2\tau = \frac{\bar{\rho} \sinh C_2\tau}{J_P C_2} = E(\tau), \end{aligned} \quad (\text{D.13g})$$

the middle inequality being (D.13b)–(D.13c) and the last step $Wz_{CoM} = J_P C_2^2$ again. **The fitted cosh lies inside the very band that contains the data**, and inside it with the same envelope $E(\tau)$, not a larger one.

The inequality in (D.13g) is an equality when $|\delta|$ takes its extreme admissible value $\bar{\rho}/\dot{M}$. So the band is not a container that happens to hold the fit: to first order it is exactly the tube swept by the admissible onset family

$$\left\{ C_1(\cosh C_2(\tau - \delta) - 1) : |\delta| \leq \bar{\rho}/\dot{M} \right\},$$

whose two extreme members *are* the band edges $\omega_{\text{nom}} \pm E$. Every curve the minimiser can return is in the band, and every part of the band is reached by some such curve.

That makes the geometric statement and the threshold statement one and the same. Since $\Delta M_{\text{crit},e} = \dot{M}\delta$, the assertions “ $|\Delta M_{\text{crit},e}| \leq \bar{\rho}$ ” and “the fitted cosh lies between $\omega_{\text{nom}} \pm \bar{\rho} \sinh(C_2\tau)/(J_P C_2)$ ” are two readings of a single inequality, one about the answer and one about the curve. Nothing beyond (D.13a) and the projection was used to pass between them.

Two qualifications. First, the pinning is doing real work here. Were C_1 free as well, the fit would be $R \cosh(C_2\tau + \psi) - C_1$ with $R \neq C_1$, and the amplitude channel is governed by $P = \int \cosh(C_2s) \rho ds$ rather than by Q ; P carries no factor that vanishes at $s = 0$ and is not bounded by $\bar{\rho}$, so the fitted curve could then leave the band. It is precisely because $C_1 = K\dot{M}$ is fixed that the only freedom left is the one E bounds. Second, (D.13g) is first order in g , and (D.13f) says what that costs. Over the operating box, with $\bar{\rho} = 12.04$ mN m and $C_2 = 5.05$:

$\dot{M}$ N m/s	$ \delta _{\text{max}} = \bar{\rho}/\dot{M}$ s	$g = C_2 \delta _{\text{max}}$	(D.13f)/ E at $x = 3$	dropped term at $\tau = 0$ % of peak rate
0.10	0.1204	0.608	1.377	2.10
0.20	0.0602	0.304	1.169	0.51
0.45	0.0268	0.135	1.071	0.10
0.80	0.0150	0.076	1.039	0.03
1.20	0.0100	0.051	1.026	0.01

The exact band is wider than E by 2.6% at the fast ramp and 38% at the slowest, and the widening is a second-order effect of an *a priori* onset displacement that is itself a fifth of the

window at $\dot{M} = 0.10$; the displacements actually measured are two orders smaller, so in the campaign $g \lesssim 10^{-3}$ and the two forms are indistinguishable. The last column is the one place the first-order form fails outright: at $\tau = 0$ the envelope E vanishes while $C_1(\cosh g - 1)$ does not, so a curve whose onset moved early is strictly outside the band there. What it is outside by is at most 2.1% of the peak nominal rate, and it is at the one end of the window where nothing is being measured. Every claim above is verified in `analysis/absorption_check.py`.

The general statement, and what it costs. Strip away every special feature — the cosh, the pinned amplitude, the hyperbolic identities — and ask the bare question. Let M be *any* set of candidate curves in L^2 of the window, let the record be $y = f + e$ with $f \in M$, and let $\hat{g} \in \arg \min_{g \in M} \|y - g\|$. Then

$$\|\hat{g} - f\| \leq \|\hat{g} - y\| + \|y - f\| \leq \|f - y\| + \|y - f\| = 2\|e\|, \quad (\text{D.14a})$$

the first step the triangle inequality and the second the definition of $\hat{g}$ together with $f \in M$ — if f is admissible, the minimiser cannot do worse than f does. That is the whole proof. It uses no linearisation, no convexity, no identifiability, no smoothness, and no property of M whatsoever. So the answer to whether the fitted curve must lie in a band around the nominal is *yes, in general and cheaply*, provided two things hold: the nominal is in the class, and the minimum is global. Both are worth naming as assumptions rather than left implicit, and both are examined below.

Three things are given up in exchange for that generality, and they are exactly the three the earlier derivation had.

The constant is 2, and it is sharp. Take $M = \{f, g\}$, two points, and $e = \frac{1}{2}(g - f) + \varepsilon$ chosen so that y falls just on g 's side. The minimiser returns g , and $\|\hat{g} - f\| = \|g - f\| \rightarrow 2\|e\|$. Nothing about (D.14a) can be improved without assuming something about M . Assuming a little recovers the factor. If M is an affine subspace $f + V$ then $\hat{g} - f = P_V e$ exactly, and orthogonal projection onto a subspace is non-expansive, so

$$\|\hat{g} - f\| = \|P_V e\| \leq \|e\|, \quad (\text{D.14b})$$

with equality iff $e \in V$. For a curved M of reach r in the sense of Federer — the largest radius within which nearest-point projection is single-valued — the projection is Lipschitz with constant $r/(r - d)$ at distance d , giving

$$\|\hat{g} - f\| \leq \frac{\|e\|}{1 - \|e\|/r} \quad (\|e\| < r), \quad (\text{D.14c})$$

which is (D.14b) as $\|e\|/r \rightarrow 0$ and (D.14a) at $\|e\| = r/2$. Our class is curved, so the honest constant is (D.14c); measured on it, r runs from 0.068 at $\dot{M} = 0.10$ to 3.35 at 1.20, against $\|E\| = 0.0174$, so the constant is 1.348 at the slowest ramp and 1.005 at the fastest. An adversarial search over 602 profiles with $|e| \leq E$ — including $e = \pm E$ and random sign-switching ones — finds a worst realised ratio of 1.0005: just above the affine value, as a curved class requires, and far inside both bounds.

It is an L^2 statement, not a pointwise one. (D.14a) bounds $\|\hat{g} - f\|$, not $|\hat{g}(\tau) - f(\tau)|$, and no amount of work on it will produce the second, because a small L^2 error is compatible with a large excursion on a short interval. The pointwise band of (D.13g) is available only because the class is one-parameter and (D.13e) says exactly what a member looks like. This is the real division of labour: (D.14a) is what guarantees closeness, and the parametrisation is what converts closeness into a band.

What the direction buys. (D.14a) uses only the *size* of e . The projection route (D.13b) uses $\langle e, \chi \rangle$ — where e points, not merely how large it is. Here that turns out to buy nothing, and for an instructive reason: $E \propto \sinh C_2 \tau$ and $\chi \propto \sinh C_2 \tau$ are parallel, so Cauchy–Schwarz is tight and the two agree identically. Numerically,

$$\frac{\langle E, |\chi| \rangle}{\|\chi\|^2} = \frac{\|E\|}{\|\chi\|} = \frac{\bar{\rho}}{\dot{M}}$$

to seven decimals at every ramp rate. So (D.13b), (D.14b) and (D.13c) are three routes to one number, and the arrangement noted at the head of this passage was not a trick — but it was not doing the work either.

Where the argument actually fails. Not in the constant, and not in the geometry, but in $f \in M$. Nothing above survives if the class does not contain the nominal; the general form is

$$\|\hat{g} - f\| \leq 2\|e\| + d_M, \quad d_M \triangleq \text{dist}(f, M), \quad (\text{D.14d})$$

by the same two lines with $\|f - y\|$ replaced by the distance to the nearest admissible curve. Since C_2 is pinned per configuration, a misidentified C_2 is precisely this failure — and it is not hypothetical, the fitted values across the campaign spanning 3.50 to 8.00. Fitting noiseless data generated at C_2^{true} with the model pinned at 5.046, at $\dot{M} = 1.20 \text{ N m/s}$:

C_2^{true}	error %	d_M	$d_M/2\ e\ $	onset bias ms	threshold bias mN m
4.100	−18.7	0.0102	0.38	+36.3	+43.5
4.600	−8.8	0.0053	0.18	+18.0	+21.6
5.046	0.0	0.0000	0.00	0.0	0.0
5.600	+11.0	0.0083	0.20	−24.5	−29.4
6.400	+26.8	0.0246	0.44	−64.2	−77.1
8.000	+58.5	0.0837	0.82	−150.0	−180.0

A C_2 wrong by 9% biases the identified threshold by 21.6 mN m with no noise present at all — $1.8\times$ the entire a priori budget $\bar{\rho} = 12.04 \text{ mN m}$ — and at 58% the fit saturates the $\pm 150 \text{ ms}$ search limit. This is a bias, not a bound: it does not average down over runs, and no error model on ρ covers it.

The other two directions. A reader who counts parameters will object that the model has four of them and this appendix has projected onto two. The objection is right and the answer is that they belong to different stages. Differentiating $g = C_1(\cosh C_2(\tau - \delta) - 1) + c$ at the nominal gives four tangent directions,

$$\begin{aligned} u &= \partial g / \partial C_1 = \cosh C_2 \tau - 1, & \chi &= \partial g / \partial \delta = -C_1 C_2 \sinh C_2 \tau, \\ \mathbf{1} &= \partial g / \partial c, & v &= \partial g / \partial C_2 = C_1 \tau \sinh C_2 \tau, \end{aligned} \quad (\text{D.15a})$$

and only the last of them leaves $\text{span}\{\cosh, \sinh, 1\}$ — it carries the factor τ , so it is the resonant direction of §4 and the one the shift identity cannot reach.

The onset sweep holds C_1 and C_2 fixed, so at that stage u and v are not available and the projection onto $\text{span}\{\chi, \mathbf{1}\}$ is exactly right. What the two frozen directions contribute is not absorption but *bias*: if the calibration supplies $C_1(1 + \varepsilon_1)$ and $C_2(1 + \varepsilon_2)$, the deviation the sweep must accommodate is not e_ω but $e_\omega - \varepsilon_1 u - \varepsilon_2 C_2 v$, and the normal equation (D.11) picks up two extra terms. Both integrate in closed form. With $x = C_2 \tau_{\text{end}}$,

$$\begin{aligned} \langle u, \chi \rangle &= -C_1^2 A(x) / 1, & \langle v, \chi \rangle &= -C_1^2 D(x) / C_2, \\ \|\chi\|^2 &= C_1^2 C_2 B(x), \end{aligned} \quad (\text{D.15b})$$

$$\begin{aligned} A(x) &= \frac{1}{4} \cosh 2x - \cosh x + \frac{3}{4}, & B(x) &= \frac{1}{4} \sinh 2x - \frac{1}{2} x, \\ D(x) &= \frac{1}{4} x \sinh 2x - \frac{1}{8} \cosh 2x + \frac{1}{8} - \frac{1}{4} x^2, \end{aligned} \quad (\text{D.15c})$$

so that, carrying each through $\Delta M = \dot{M} \delta$ and $C_1 W z_{C_0 M} = \dot{M}$,

$$\Delta M_{\text{crit}, \varepsilon} = \varepsilon_1 \frac{\dot{M}}{C_2} \frac{A(x)}{B(x)} + \varepsilon_2 \frac{\dot{M}}{C_2} \frac{D(x)}{B(x)}. \quad (\text{D.15d})$$

Both scale with $\dot{M}/C_2$, the moment swept in one plant time constant, and D/B exceeds A/B by a factor 2.4 to 3.6 over the operating windows: *a misidentified exponent is about three times as costly as a misidentified amplitude*. At $x = 3.735$ and $\dot{M} = 0.45$ N m/s, (D.15d) gives 0.82 and 2.90 mN m per 1% against 1.30 and 3.72 measured on the full nonlinear fit — the gap being the linearisation and the pre-onset segment, which the closed form does not carry. A and D agree with direct quadrature to machine precision.

Two things keep this from entering the reported budget. The calibration is per configuration and carries no direction split, so ε_1 and ε_2 are common to the two runs and (D.15d) is antisymmetric in $\dot{M}$; measured on the full nonlinear tip-over the half-sum rejects 99.77 to 99.98% of it, leaving 0.000 mm at $\pm 10\%$ amplitude and $\pm 5\%$ exponent error. And what does survive goes to the *weight* of (34), which reads the difference rather than the sum, at 0.027% of W per 1% of amplitude error and 0.077% per 1% of exponent error. Verified in `analysis/two_sided_cancellation.py`.

That is the useful conclusion of this passage, and it is not the one the band statement suggests. The band argument is safe but nearly empty; the sensitivity that matters sits in the assumption it takes for granted. It also says what the per-configuration PNLS calibration is *for*: not to improve the fit cosmetically, but to buy $f \in M$, which is the hypothesis every bound in this appendix rests on. Verified in `analysis/general_band.py`.

step	what changes	why it is allowed
1–2	$e_\omega \rightarrow$ Duhamel double integral	(D.2), the exact solution
3	double integral $\rightarrow \int \rho w$	Fubini on the triangle
4	—	$w \geq 0$, $\int w = 1$, $w \downarrow$
5	$\int \rho w \rightarrow \rho_{\max}$	(i) and (ii)
6	$\rho_{\max} \rightarrow \bar{\rho}$	(iii) plus Chebyshev (D.4)
7	$\bar{\rho} \rightarrow \rho_\varphi/7 + \rho_{GE}/5$	(95)

Why the constants differ from the left half. The left half asks how far the rate curve is displaced, so the kernel integral Ψ multiplies $\bar{\rho}$ and the products ΨR climb to $\frac{1}{2}$ and 1. The right half asks how far the *minimiser* moves, and the minimiser is found by correlating against $\sinh(C_2\tau)$ — a normalised operation, $\int w = 1$, with no kernel gain left over. So there $\bar{\rho}$ enters bare, and the constants are the R ’s themselves, $\frac{1}{7}$ and $\frac{1}{5}$. The threshold bound is the stronger of the two by roughly a factor of four, which is why it, and not the relative rate bound, is the number carried into the offset budget.

7. The threshold bound by hand

The manuscript reports only the threshold bound, and that branch uses neither Ψ nor x : the weight of (D.13) is normalised, so $\bar{\rho}$ enters bare and the constants are the R ’s themselves. What is left is arithmetic. No hyperbolic function appears below, C_2 never enters, and the window length is never solved for — $\frac{1}{7}$ and $\frac{1}{5}$ are suprema over *all* x , so no x has to be known to use them.

Given the weight W , the height z_{CoM} , the arm $a = l_p + p_{\text{off}}$, the mass m , a CAD inertia J_{CoM} , the ramp rate $\dot{M}$, the tilt cap $\varphi_{\max}$ and the ground-effect slope β_M , six lines:

$$(1) \quad J_P \leq J_{CoM} + m(z_{CoM}^2 + a^2) \quad \text{— (82) at its upper end}$$

$$(2) \quad \Delta M_{\text{win}} \leq (6\varphi_{\max} J_P \dot{M}^2)^{1/3} \quad \text{— (93)}$$

$$(3) \quad \rho_{\varphi, \max} = \frac{1}{2} W a \varphi_{\max}^2 \quad \text{— the gravity remainder at the cap}$$

$$(4) \quad \rho_{GE, \max} = |\beta_M| \Delta M_{\text{win}} \varphi_{\max}$$

$$(5) \quad |\Delta M_{\text{crit}}| \leq \rho_{\varphi, \text{max}}/7 + \rho_{GE, \text{max}}/5 \quad \text{--- (97b)}$$

(6) divide by W_{min} for the equivalent offset.

Over the geometric box of Sec. VI-D, with $W = 31.59$ N, $z_{CoM} = 0.30$ m, $m = 3.221$ kg, $J_{CoM} = 0.0537$ kg m², $\varphi_{\text{max}} = 10^\circ$, $\beta_M = 0.0345$ rad⁻¹ and $W_{\text{min}} = 30.08$ N:

axis	$\dot{M}$ N m/s	J_P kg m ²	ΔM_{win} N m	$\rho_{\varphi, \text{max}}$ mN m	$\rho_{GE, \text{max}}$ mN m	$ \Delta M_{\text{crit}} $ mN m	offset mm
roll ($a = 0.160$)	0.10	0.4260	0.165	76.98	0.99	11.20	0.372
	0.45	0.4260	0.449	76.98	2.70	11.54	0.384
	1.20	0.4260	0.863	76.98	5.20	12.04	0.400
pitch ($a = 0.130$)	0.10	0.3980	0.161	62.55	0.97	9.13	0.304
	0.45	0.3980	0.439	62.55	2.64	9.46	0.315
	1.20	0.3980	0.844	62.55	5.08	9.95	0.331

Three readings.

The inertia hardly matters. J_P enters only through ΔM_{win} , which enters only through $\rho_{GE, \text{max}}$, which is 2–10% of the total; and it enters as a cube root. Halving J_P moves the bound by -1.8% , doubling it by $+2.2\%$. A crude inertia is therefore good enough, which is worth knowing because (82) is the one input here that is not directly measured.

Line (5) is where the work went. Properties 1 and 2 of the weight alone — $w \geq 0$ and $\int w = 1$ — already give $|\Delta M_{\text{crit}}| \leq \rho_{\text{max}} = \rho_{\varphi, \text{max}} + \rho_{GE, \text{max}}$, without Chebyshev and without any R . That is 2.70 mm at worst over the box, which *exceeds* the 1.64 mm validation RMS: a guarantee weaker than the measured performance, and so of no use. The factor of seven that $\frac{1}{7}$ and $\frac{1}{5}$ supply is what brings it to 0.40 mm and clears the measurement by $4.2\times$.

The rate bound is not needed for this. The same construction also bounds the relative rate deviation, with $\frac{1}{2}$ and 1 in place of $\frac{1}{7}$ and $\frac{1}{5}$ and a factor $\Psi = x \coth(x/2)$; that is Secs. 4–5, it is weaker by about a factor of four, and no number above depends on it.

8. A worked example, checkable by hand

On which x to use. Two values circulate below and they must not be confused. The cubic shortcut (92) supplies $\bar{x} = (6\Lambda_{\text{max}})^{1/3}$, an *upper bound* on x obtained by replacing $\Lambda(x)$ with its leading term $x^3/6$. The exact value solves

$$\Lambda(x) = \Lambda_{\text{max}} \triangleq \frac{\varphi_{\text{max}} W z_{CoM} C_2}{\dot{M}}, \quad (\text{D.14})$$

which has a unique root because Λ increases strictly from 0. The shortcut exists only to spare (97b) a transcendental solve it does not need — there x appears nowhere, and ΔM_{win} enters through a channel worth under a tenth of the total. Wherever x itself appears, (D.14) is solved instead, for two reasons. It is loose: at the slowest ramp Λ exceeds $x^3/6$ by $3.6\times$, and anything downstream of $\sinh$ exponentiates that. And it is not automatically valid: substituting an upper bound into $R(x) \sinh x$ presumes that product monotone, which is true but is not among the facts Sec. 5 proves, whereas solving (D.14) presumes nothing.

Both reduction factors close in finitely many terms, so the rows below need no quadrature:

$$\begin{aligned} R_{\varphi}(x) &= \frac{A_{\varphi}(x)}{x\Lambda(x)^2}, & A_{\varphi}(x) &= \frac{1}{4} \sinh 2x - \frac{1}{2}x - 2x \cosh x + 2 \sinh x + \frac{1}{3}x^3, \\ R_{GE}(x) &= \frac{A_{GE}(x)}{x^2\Lambda(x)}, & A_{GE}(x) &= x \cosh x - \sinh x - \frac{1}{3}x^3, \end{aligned} \quad (\text{D.15})$$

each being $\int_0^x \Lambda^2$ and $\int_0^x u\Lambda$ integrated by parts against $\int \sinh^2 = \frac{1}{4} \sinh 2x - \frac{1}{2}x$ and $\int u \sinh u = u \cosh u - \sinh u$.

Take the roll axis at the fastest ramp, evaluated over the geometric box of Sec. VI-D rather than at any fitted constant: $W = 31.59$ N (fully ballasted, $m = 3.221$ kg), $z_{CoM} = 0.30$ m, $a = l_p + p_{\text{off}} = 0.160$ m, $\dot{M} = 1.20$ N m/s, $\varphi_{\text{max}} = 10^\circ = 0.17453$ rad, $\beta_M = 0.0345$ rad $^{-1}$, $J_{CoM}^{\text{CAD}} = 0.0537$ kg m 2 .

step	expression	value
geometry	$Wz_{CoM} = Wz_{CoM}$	9.477 N m
(82) lower	$J_P \geq m(z_{CoM}^2 + a^2)$	0.3724 kg m 2
(82) upper	$J_P \leq J_{CoM}^{\text{CAD}} + m(z_{CoM}^2 + a^2)$	0.4261 kg m 2
	$C_2 \leq \sqrt{gz_{CoM}/(z_{CoM}^2 + a^2)}$	5.045 s $^{-1}$
tilt cap $\rightarrow \Lambda$	$\Lambda_{\text{max}} = \varphi_{\text{max}} Wz_{CoM} C_2 / \dot{M}$	6.954
(D.14)	solve $\sinh x - x = 6.954$	$x = 2.9930$
	check: $\sinh(2.9930) - 2.9930$	6.9548
(92)	$\bar{x} = (6\Lambda_{\text{max}})^{1/3}$, the shortcut	3.4685
(93)	$\tau_{\text{end}} \leq (6\varphi_{\text{max}} J_P / \dot{M})^{1/3}$	0.719 s
(93)	$\Delta M_{\text{win}} = \dot{M} \tau_{\text{end}}$	0.863 N m
channel	$\rho_{\varphi, \text{max}} = \frac{1}{2} W a \varphi_{\text{max}}^2$	76.98 mN m
(93)	$\rho_{GE, \text{max}} = \beta_M \Delta M_{\text{win}} \varphi_{\text{max}}$	5.20 mN m
(D.15)	$\Lambda(x) = 6.9548, \Lambda(x)^2$	48.3690
(D.15)	$A_\varphi(x)$	17.2171
(D.6)	$R_\varphi(x) = 17.2171 / [(2.9930)(48.3690)]$	0.1189
(D.15)	$A_{GE}(x)$	11.0389
(D.7)	$R_{GE}(x) = 11.0389 / [(2.9930)^2 (6.9548)]$	0.1772
(95)	$\bar{\rho} = R_\varphi \rho_{\varphi, \text{max}} + R_{GE} \rho_{GE, \text{max}}$ (against $\rho_{\text{max}} = 82.2$ mN m)	10.08 mN m
(96)	$\Psi(x) = x \coth(x/2)$	3.309
(96)	$\Psi \bar{\rho} / \Delta M_{\text{win}}$	3.86%
(97)	$(\frac{1}{2} \rho_{\varphi, \text{max}} + \rho_{GE, \text{max}}) / \Delta M_{\text{win}}$	5.06%
(97b)	$\rho_{\varphi, \text{max}} / 7 + \rho_{GE, \text{max}} / 5$ $\div W_{\text{min}}$	12.04 mN m 0.400 mm
(D.3)	envelope only, $\Psi \rho_{\text{max}} / \Delta M_{\text{win}}$	31.5%

Read the rows in two groups. Everything from (D.15) down to (96) is evaluated at the exact x of (D.14), so each is the value the run actually realises; the $\bar{x}$ row is carried only to show what the shortcut costs, and using it would have given $\bar{\rho} = 9.54$ instead of 10.08 mN m — an *understatement*, since R decreases. The (97) and (97b) rows are of a different kind again: they contain no x at all, so they hold whatever the window turns out to be. That is why (97b), and not the 3.86%, is what the offset budget inherits.

Two readings. First, (97) costs about a third over (96) — the price of dropping x from the statement — whereas the envelope form costs a factor of eight. Second, this is the *worst admissible* case: the box is evaluated at the 10° tilt cap, at the top of the z_{CoM} range, and with the inertia at whichever end of (82) is unfavourable for the quantity being bounded. Over the runs actually flown the excursions are $5\text{--}9^\circ$ and the exact per-run propagation, which convolves the measured ρ instead of any envelope, gives 0.04 mN m in the median and 0.15 at worst — 0.001 to 0.004 mm of offset.

9. Summary of the chain

	step	what it buys
(D.2)	$e_\omega = \int \dot{h}_\varphi(\tau_{\text{end}} - s)\rho(s)ds$	the exact deviation
(D.3)	$ \rho \leq \rho_{\text{max}}$, integrate the kernel	(90); too weak by $\sim 8\times$
(D.4)	Chebyshev, $\rho \uparrow$ against $\dot{h}_\varphi(\tau_{\text{end}} - \cdot) \downarrow$	$\rho_{\text{max}} \rightarrow \bar{\rho}$
(D.6–7)	known shapes Λ^2 and $u\Lambda$	the ratios $\bar{\rho}/\rho_{\text{max}}$
(D.6a)	$R = \int_0^1 f(xt)/f(x)dt$, $n_{\text{eff}} \uparrow$	(95): $R_\varphi \leq \frac{1}{7}$, $R_{GE} \leq \frac{1}{5}$
(96)	divide by $\omega_{\text{nom, end}}$, use $\Delta M_{\text{win}} = \dot{M}x/C_2$	$\Psi = x \coth(x/2)$; Wz_{COM} , J_P gone
	$\Psi R_\varphi \leq \frac{1}{2}$, $\Psi R_{GE} \leq 1$	(97) left half; x gone
(D.12–13)	linearise the onset, Fubini, $w \geq 0$, $\int w = 1$	$\Delta M_{\text{crit}} = -\langle \rho \rangle_w$
	Chebyshev again, $\rho \uparrow$ against $w \downarrow$	(97) right half: $\frac{1}{7}, \frac{1}{5}$

Chebyshev is used twice, for the same reason both times: the forcing is concentrated at the end of the window and every kernel in the problem weights the end of the window least. The two uses are independent — the first converts a kernel integral, the second a normalised correlation — which is why the constants they leave behind differ.